PPgEEC
Machine Learning

Fundamentals of ML and Decision Trees

Ivanovitch Silva
ivanovitch.silva@ufrn.br

Table of Contents

"A reproducible Pipeline is all you need"

What is Machine Learning?

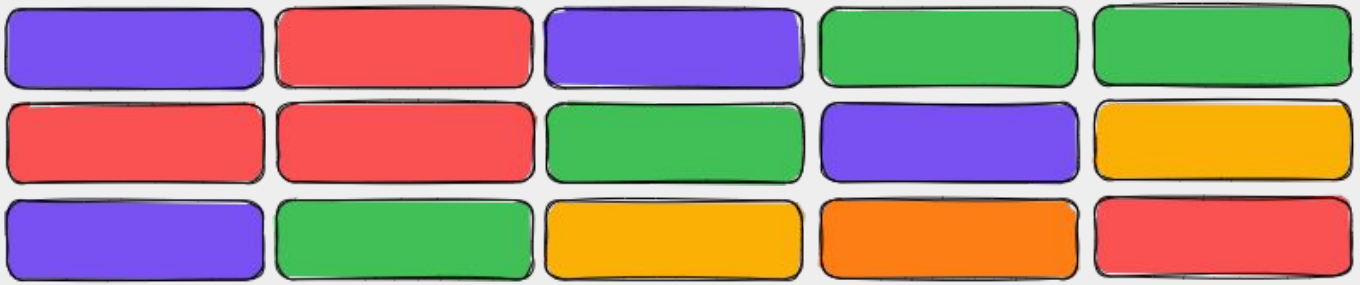

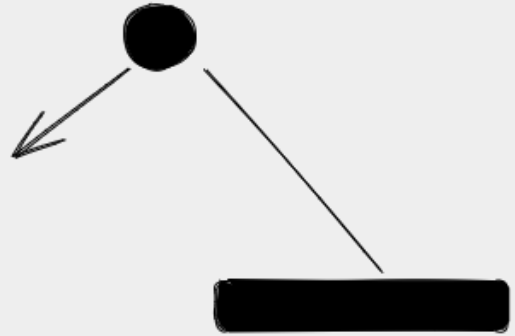

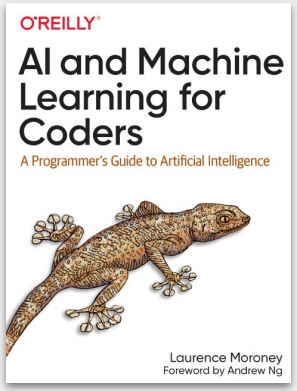


```
if (ball.collide(brick)){
    removeBrick();
    ball.dx = 1.1*(ball.dx);
    ball.dy = -1*(ball.dy);
```

Limitations of traditional programming

<activity detection>

```
if (speed < 4){
    status = WALKING;
}
```

```
if (speed < 4){
    status = WALKING;
} else {
    status = RUNNING;
}
```

```
if (speed < 4){
    status = WALKING;
} else if (speed < 12) {
    status = RUNNING;
} else {
    status = BIKING;
}
```

```
// ????
```

From coding to ML

<gathering and label data>

```
0101111000011110101
1110101010101111000
1110101010101010100
0000000000111001111
```

Label = WALKING

```
010111100011101110
000111010101011011
000111101010111000
000000000000001111
```

Label = RUNNING

```
111101110010101011
110101110101011011
111110101110010101
000111111001101111
```

Label = BIKING

```
100000000010101011
111111111000011001
000000011100111101
111111111100000001
```

Label = GOLFING

From programming to learning

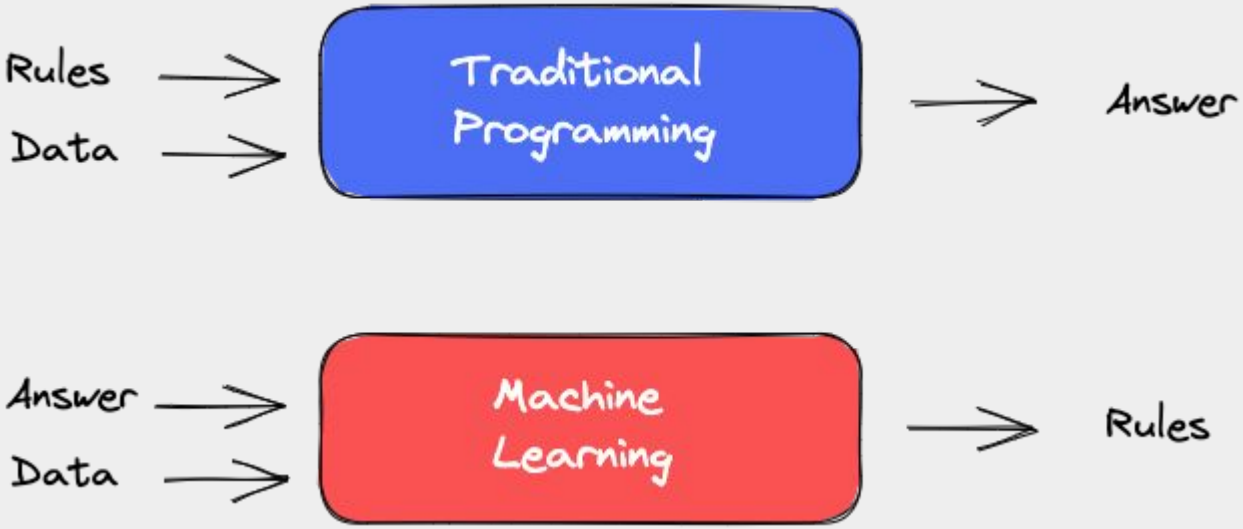

What is Machine Learning?

Machine Learning (ML): a subset of AI that often uses statistical techniques to give machines the ability to "learn" from data without begging explicitly given the instructions for how to do so. This process is known as "training" a "model" using a learning "algorithm" that progressively improves models performance on a specific task.

Computer Vision

Semantic Segmentation

203 benchmarks
2300 papers with code

Image Classification

279 benchmarks
1989 papers with code

Object Detection

264 benchmarks
1737 papers with code

Image Generation

169 benchmarks
771 papers with code

Denoising

100 benchmarks
739 papers with code

Time Series

Time Series

2 benchmarks
1127 papers with code

EEG

8 benchmarks
177 papers with code

Imputation

10 benchmarks
160 papers with code

Natural Language Processing

Language Modelling

27 benchmarks
1513 papers with code

Machine Translation

73 benchmarks
1366 papers with code

Question Answering

103 benchmarks
1307 papers with code

Sentiment Analysis

69 benchmarks
836 papers with code

Text Generation

84 benchmarks
649 papers with code

Speech

Speech Recognition

121 benchmarks
575 papers with code

Speech Synthesis

3 benchmarks
142 papers with code

Dialogue Generation

10 benchmarks
108 papers with code

Medical

Medical Image Segmentation

86 benchmarks
244 papers with code

Drug Discovery

14 benchmarks
151 papers with code

Lesion Segmentation

6 benchmarks
104 papers with code

Brain Tumor Segmentation

10 benchmarks
69 papers with code

COVID-19 Diagnosis

4 benchmarks
59 papers with code

Playing Games

Continuous Control

76 benchmarks
242 papers with code

Atari Games

65 benchmarks
213 papers with code

OpenAI Gym

9 benchmarks
112 papers with code

Graphs

Link Prediction

69 benchmarks
463 papers with code

Node Classification

77 benchmarks
370 papers with code

Graph Embedding

2 benchmarks
252 papers with code

Graph Classification

54 benchmarks
209 papers with code

Community Detection

11 benchmarks
156 papers with code

Music

Music Generation

60 papers with code

Music Information Retrieval

55 papers with code

Music Source Separation

3 benchmarks
31 papers with code

https://paperswithcode.com/sota

Machine Learning Types

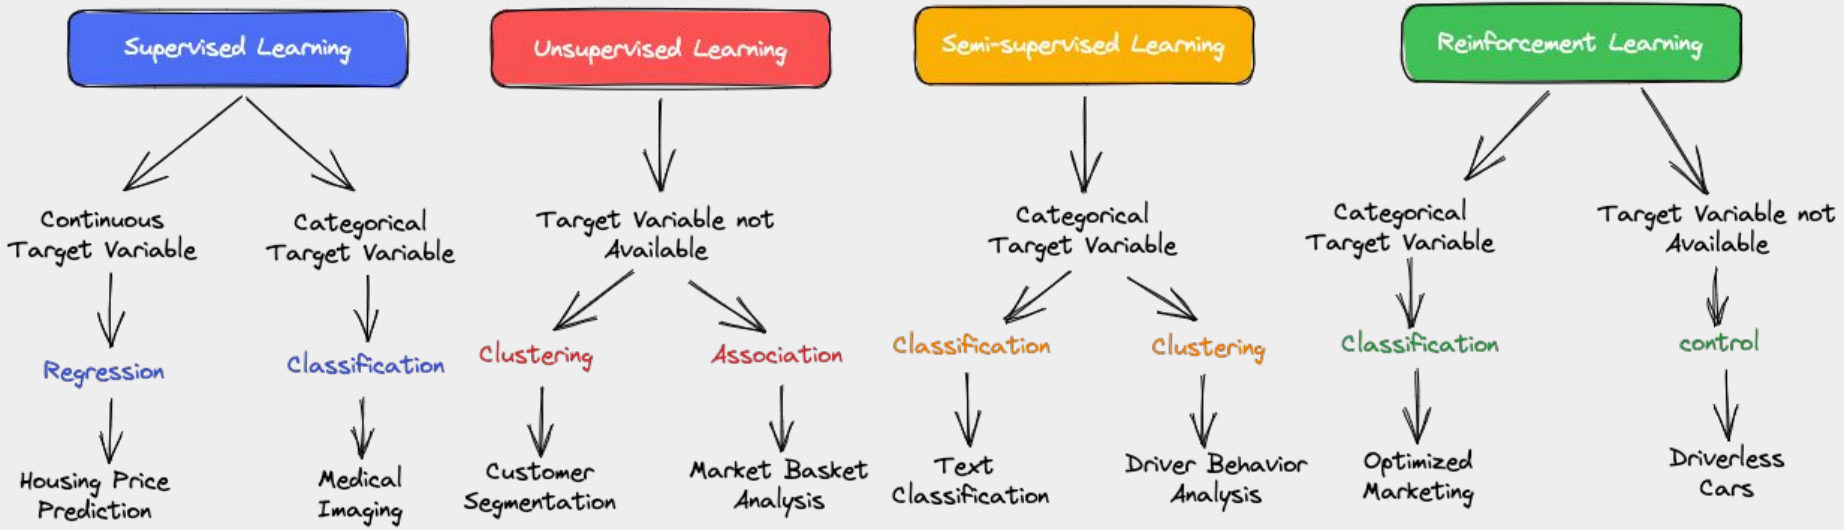

Supervised Learning

Classification Problem

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease
Yes	Yes	205	Yes
No	Yes	180	Yes
Yes	No	210	Yes
Yes	Yes	167	Yes
No	Yes	156	No
No	Yes	125	No
Yes	No	168	No
Yes	Yes	172	No

Chest Pain	Blocked Arteries	Patient Weight
Yes	No	200

Supervised Learning

Regression Problem

Height (m)	Favorite Color	Gender	Weight (kg)
1.6	Blue	Male	88
1.6	Green	Famale	76
1.5	Blue	Female	56
1.8	Red	Male	73
1.5	Green	Male	77
1.4	Blue	Female	57

Height (m)	Favorite Color	Gender
1.83	Yellow	Male

K-Neighrest Neighbors (KNN)

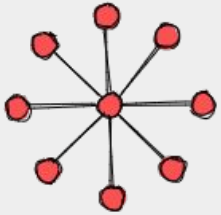

Linear Regression

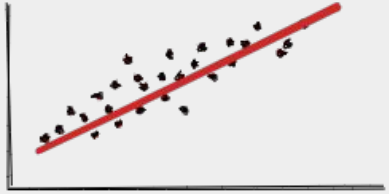

Logistic Regression

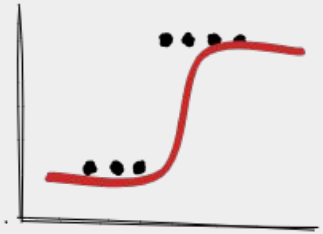

Support Vector Machines (SVM)

Decision Trees

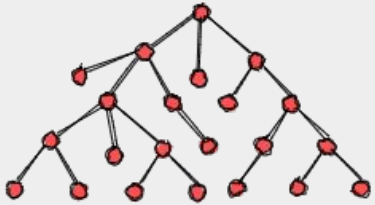

Ensemble

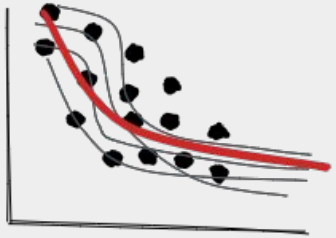

Neural Networks

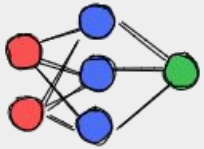

Deep Learning

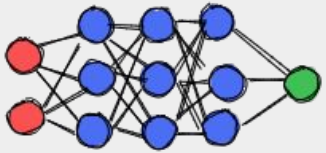

Bosting

Main Challenges Of Machine Learning

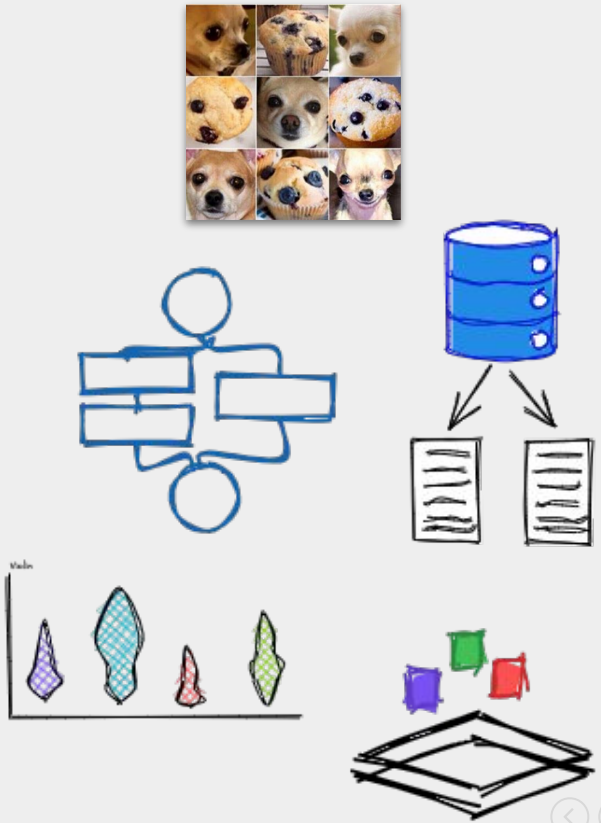

Titanic: Machine Learning from Disaster

Survived	Pclass	Name	Sex	Age	Ticket	Cabin	Embarked
0	3	Braund, Mr. Owen	Male	22	A/5 21171	NaN	S
1	1	Cummings, Mrs John	Female	38	PC 17599	C85	C
1	3	Heikkinen, Ms Laina	Female	26	STON/02	NaN	S
1	1	Futrelle, Mrs Jacques	Female	35	113803	C123	S
0	3	Allen, Mr. William	Male	35	373450	NaN	S

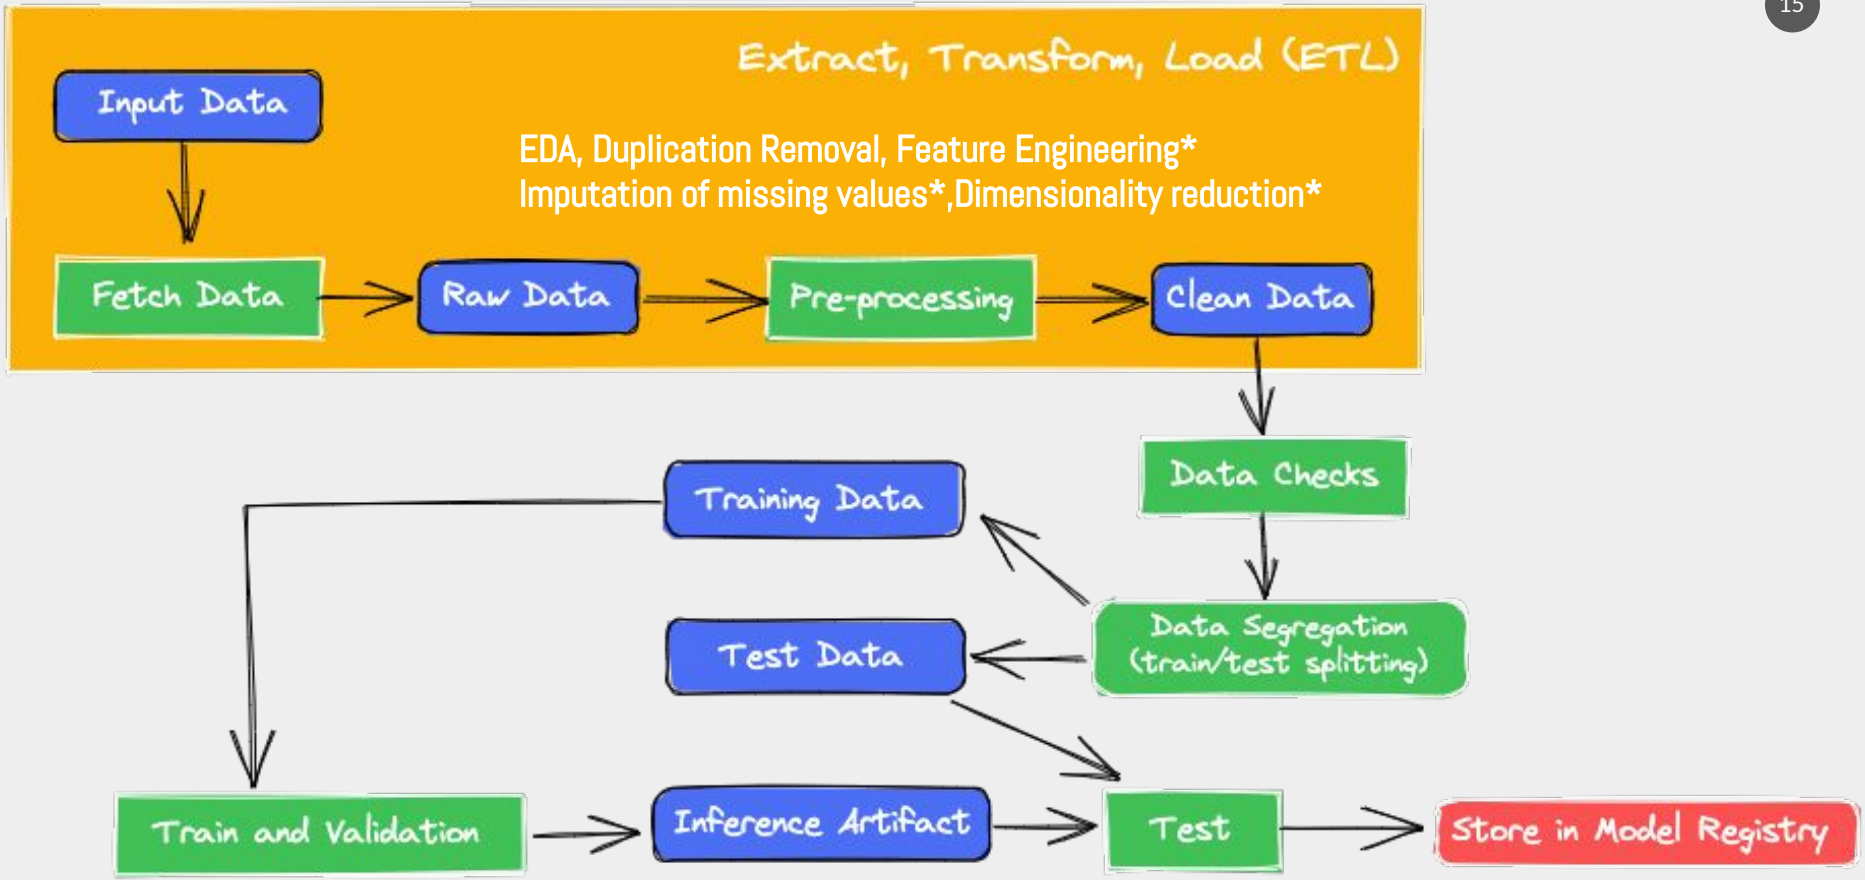

Feature Store, Categorical encoding missing values imputation, Dimensionality Reduction

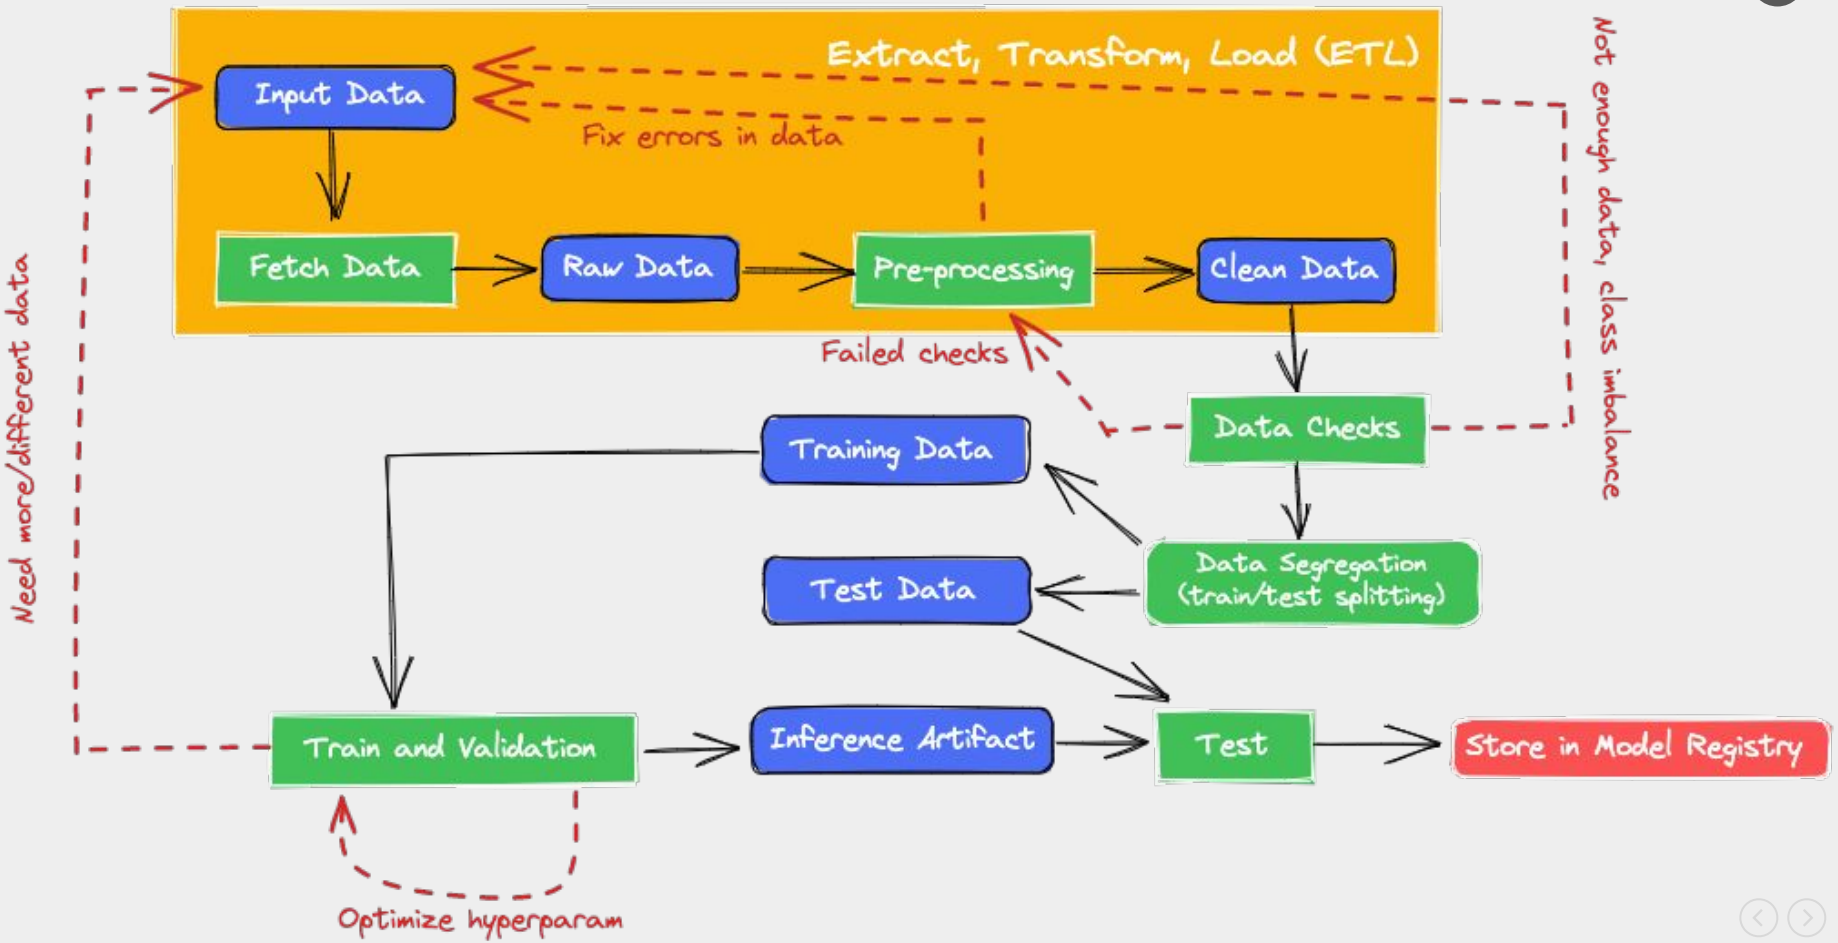
Extract, Transform, Load (ETL)
Input Data
Fix errors in data
Fetch Data
Raw Data
Pre-processing
Clean Data
Failed checks
Data Checks
Not enough data, class imbalance
Need more/different data
Training Data
Test Data
Data Segregation (train/test splitting)
Train and Validation
Inference Artifact
Test
Store in Model Registry
Optimize hyperparam

Train and Evaluate

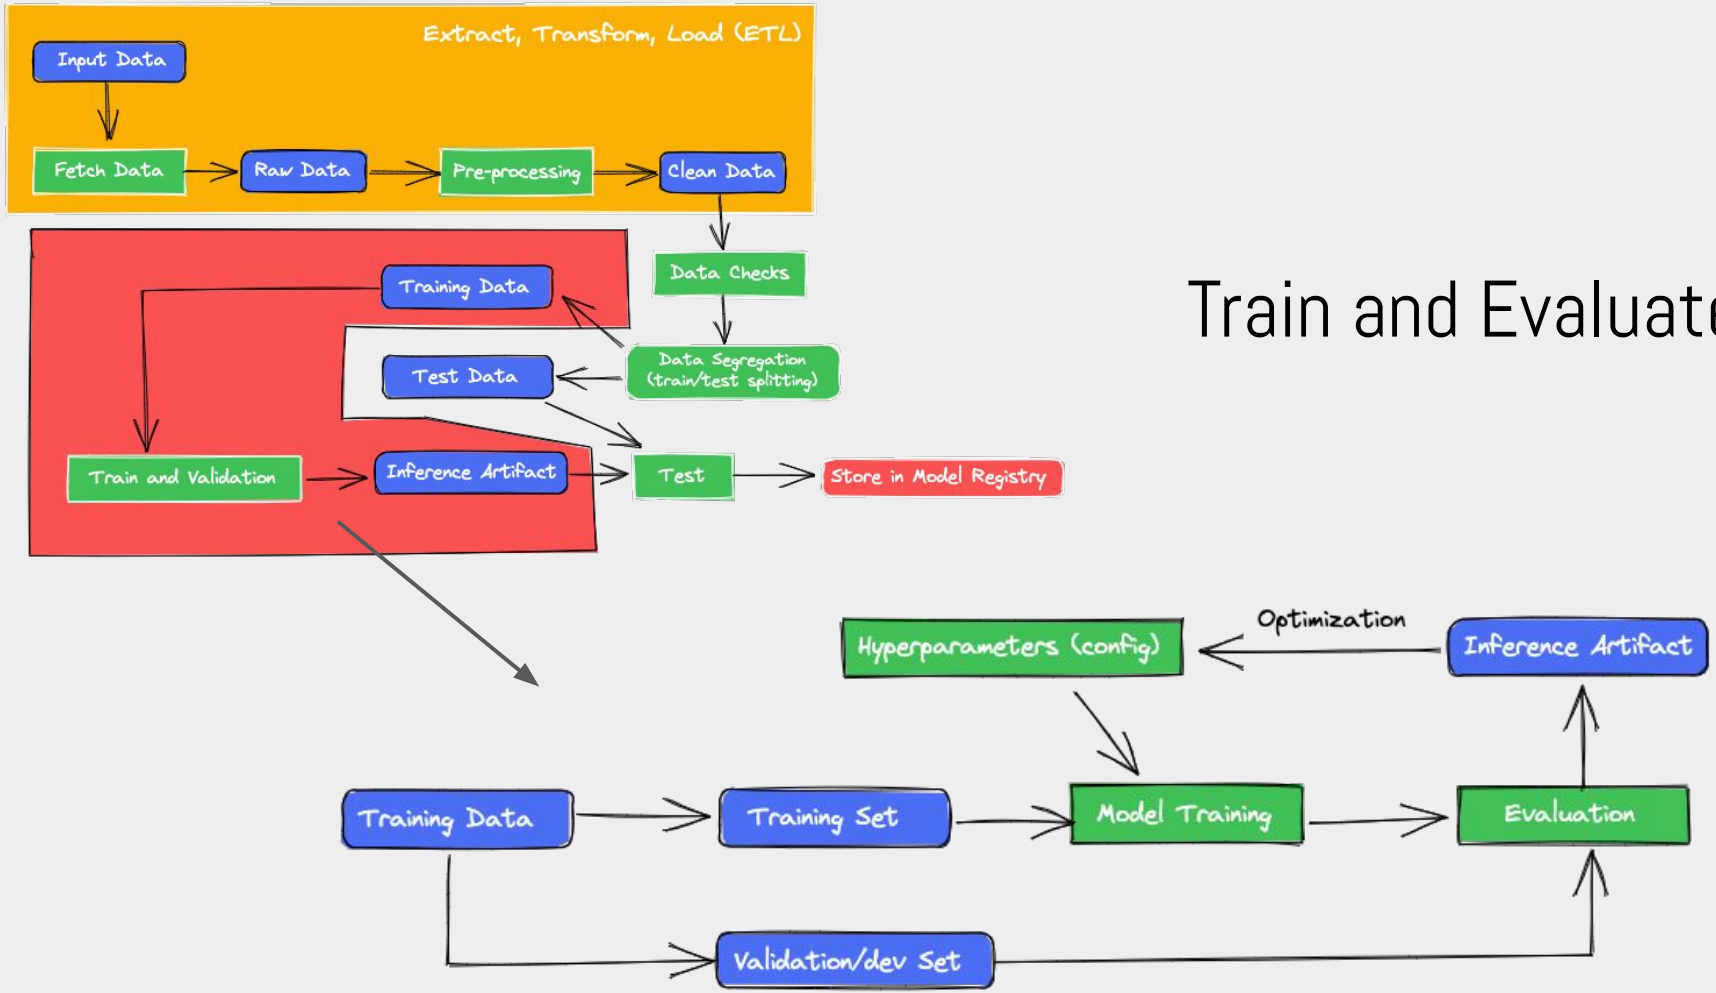

Controlled Chaos

Assume you are going to iterate A LOT

Nothing is lost
You learn something with every experiment

Give yourself time within the project deadlines

Perfection is the enemy of good
Be clear on your objective and stop once you reach it

Be systematic
Normaly, change one thing at the time

Nothing is fixed
data, code and hyperparameters

Train - Dev - Test Sets

Making good choices in how you set up your training, development, and test sets can make a huge difference in helping you quickly find a good high performance neural network.

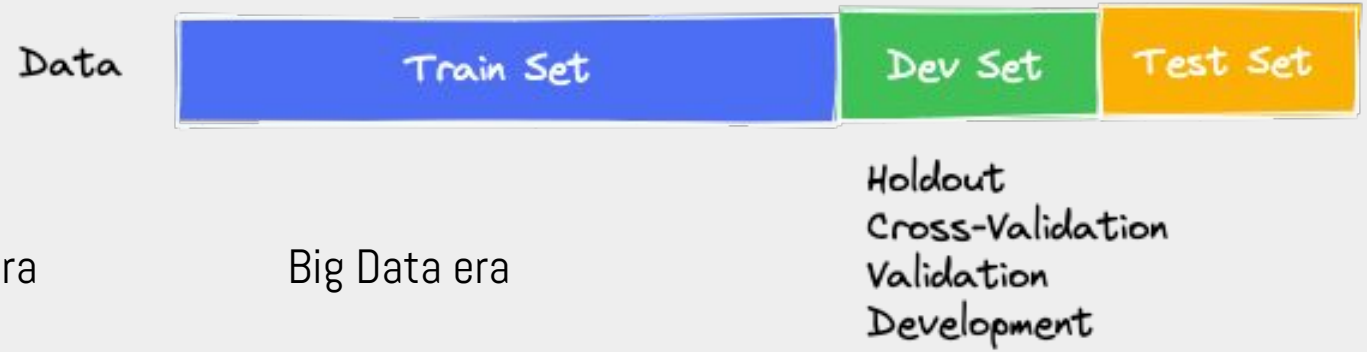

Previous ML era

- 70/30
- 60/20/20

Big Data era

- 98/1/1
- 99.5/0.25/0.25
- 99.5/0.4/0.1

Mismatched train/test distribution

Scenario: say you are building a cat-image classifier application that determines if an image is of a cat or not. The application is intended for users in rural areas who can take pictures of animals by their mobile devices for the application to classify the animals for them.

Scraped from Web Pages
100k images

Collected from Mobile Devices
<<target distribution>>
8k images

A possible option: shuffling the data

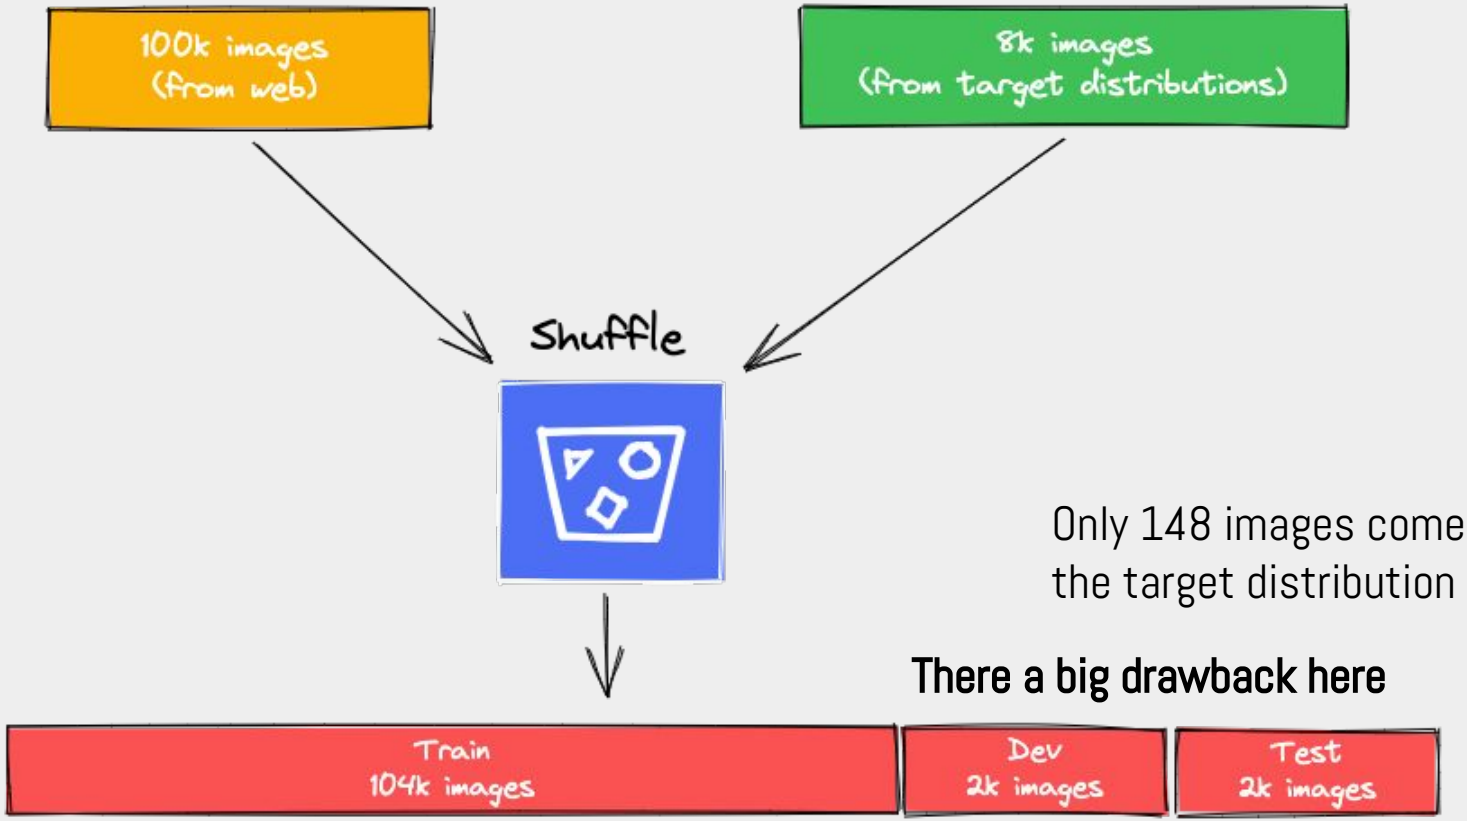

Only 148 images come from the target distribution

There a big drawback here

A better option

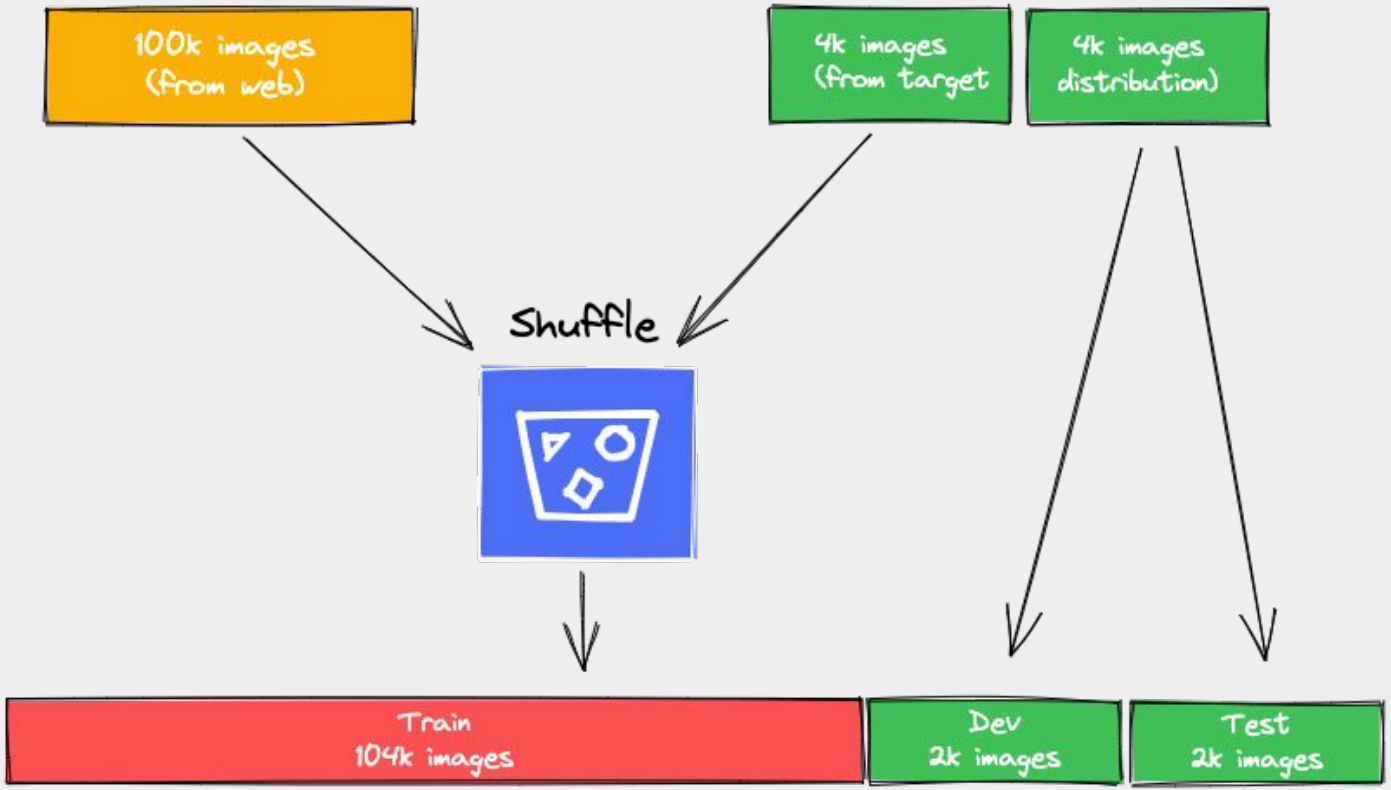

Rule of the thumb

>> make sure that the dev and test sets come from the same distribution

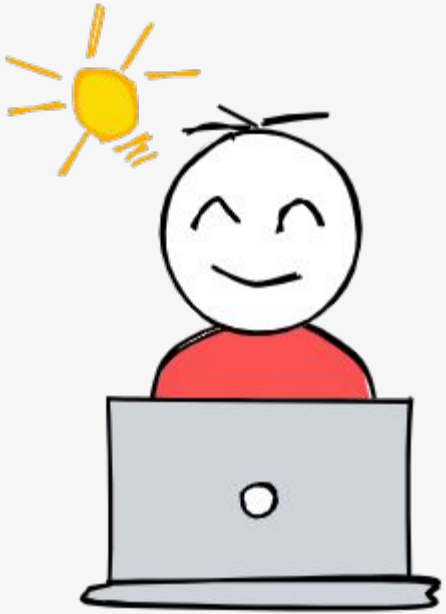

Not having a test set might be okay. (Only dev set)

Bias vs Variance

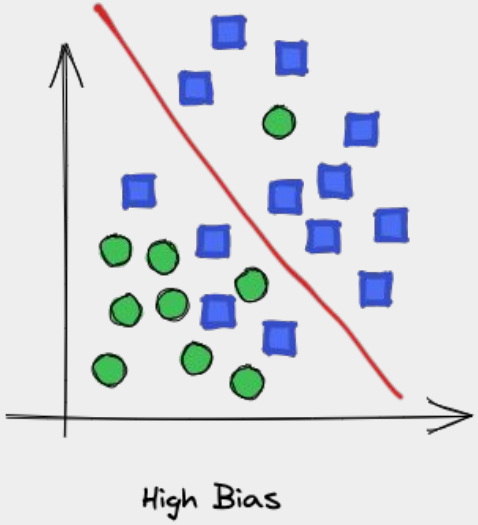

Underfitting

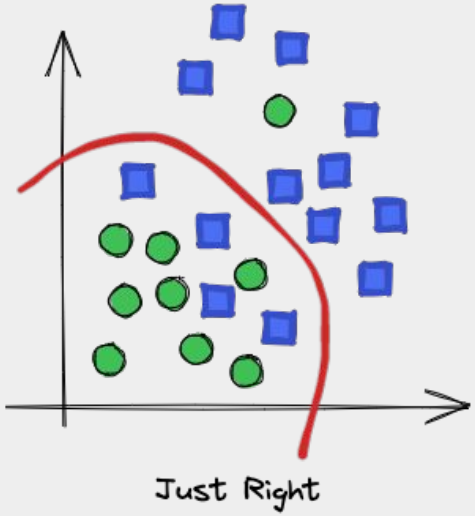

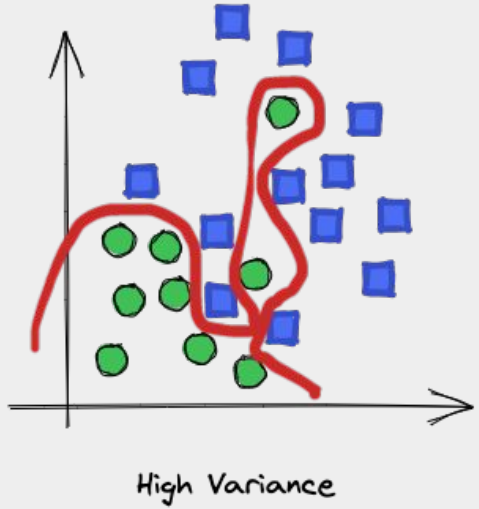

Overfitting

Bias vs Variance

Cat Classification

	Scenario #01	Scenario #02	Scenario #03	Scenario #04
Train Set Error	1%	15%	15%	0.5%
Dev Set Error	16%	16%	30%	1%
	Low Bias High Variance	High Bias Low Variance	High Bias High Variance	Low Bias Low Variance

Basic Recipe for Machine Learning

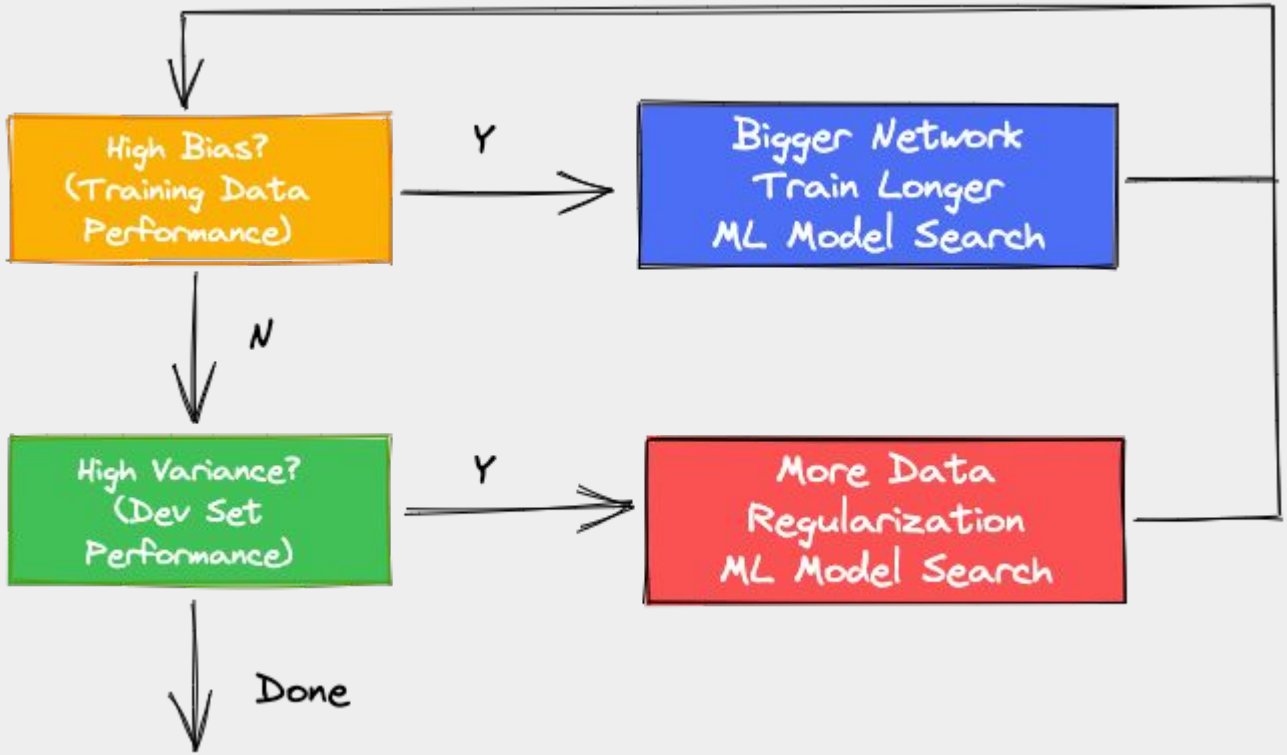

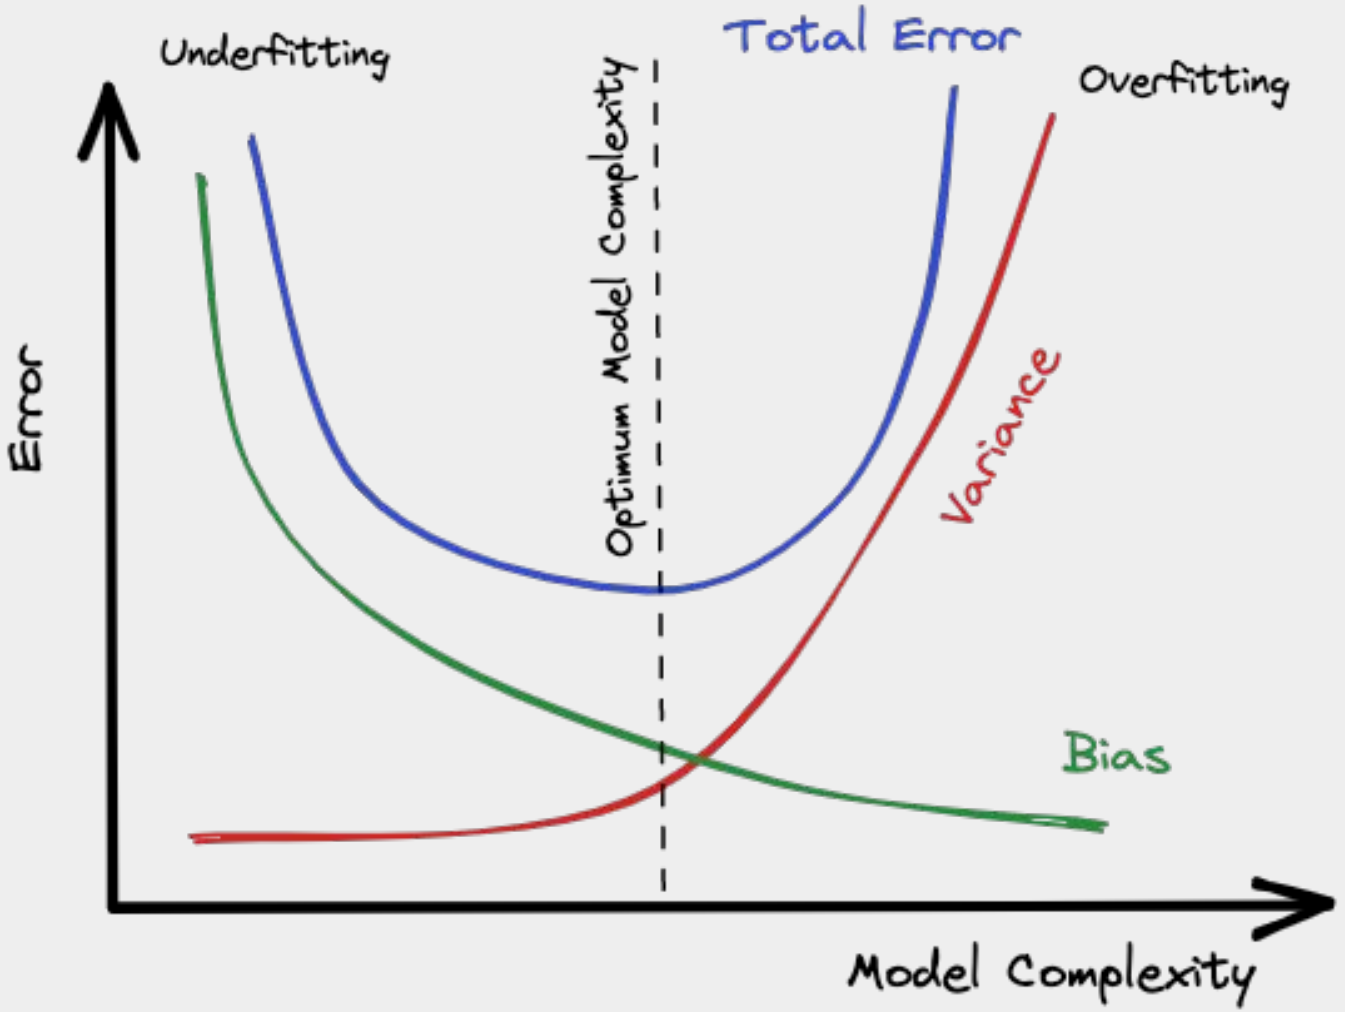
Underfitting
Total Error
Overfitting
Optimum Model Complexity
Error
Variance
Bias
Model Complexity

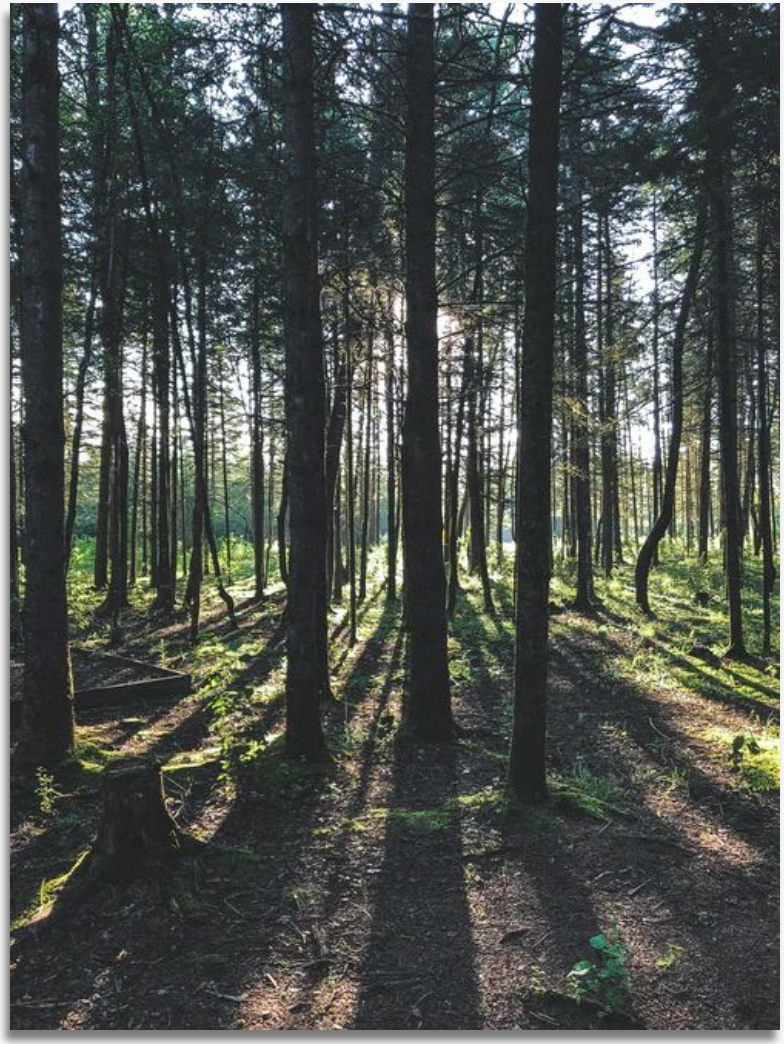

Decision Trees

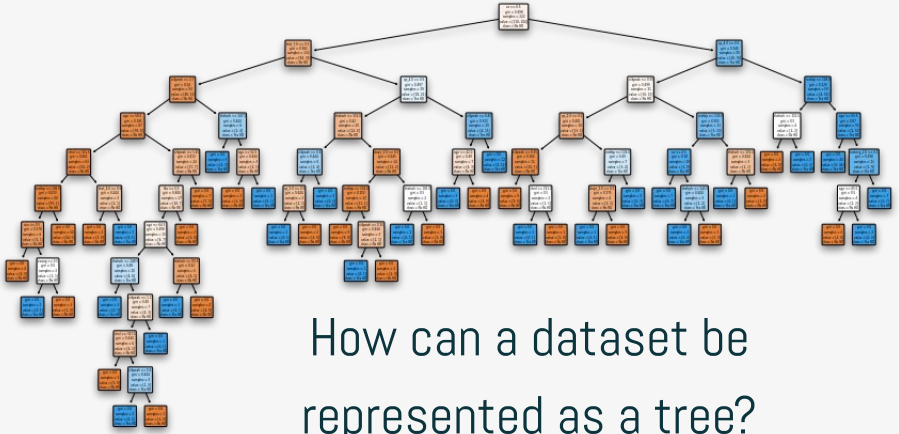

How can a dataset be represented as a tree?

Age	Marital-Status	Educational	Sex	High-Income
28	Never-Married	Bachelors	Female	<=50k
46	Never-Married	Assoc-acdm	Female	<=50k
35	Married-civ-spouse	Some-college	Male	<=50k
27	Married-civ-spouse	Bachelors	Male	>50k
59	Divorced	Some-college	Female	<=50k

Decision Tree (classification)

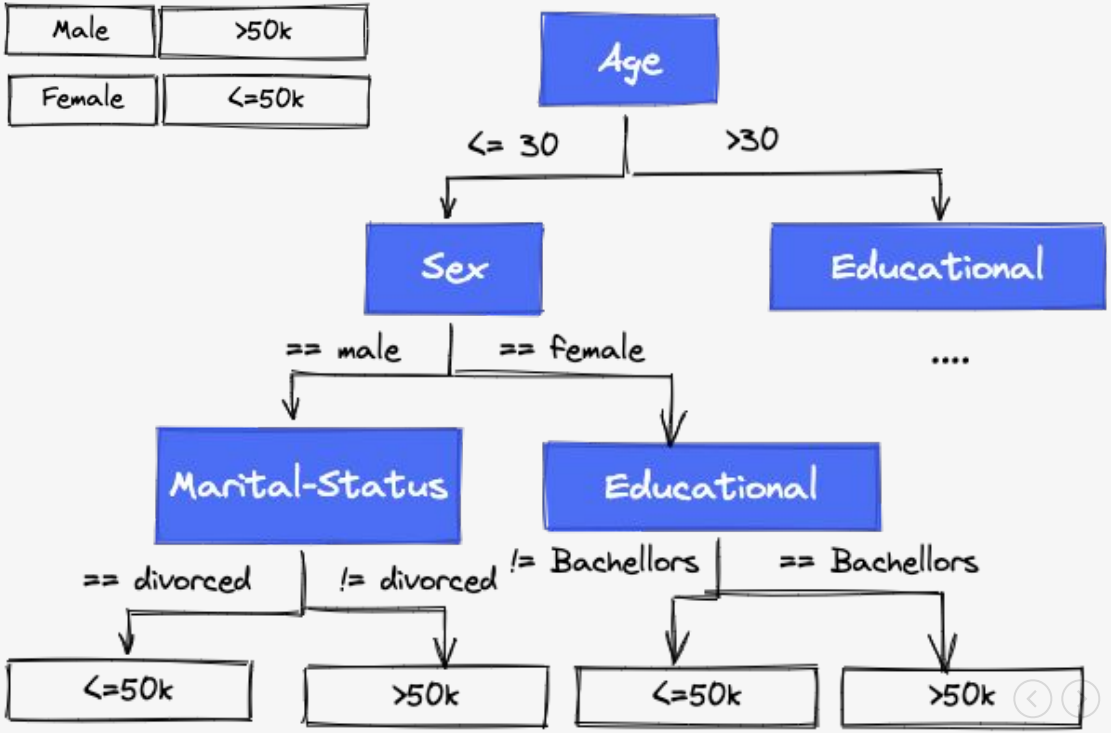

How can we split the tree?

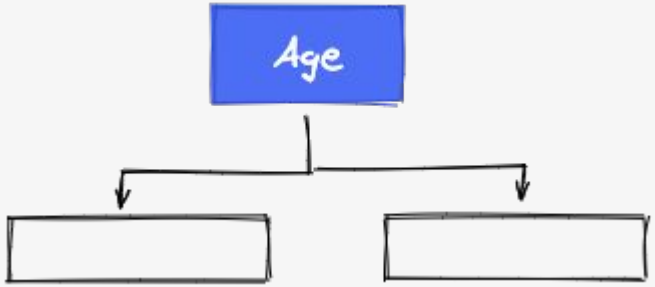

Algorithm used in Decision Trees

1. ID3 (Entropy)
2. Gini Index
3. Chi-Square
4. Reduction in Variance
    a. C4.5, pruning
5. ...

Entropy is an indicator of how messy your data is.

Why Entropy in Decision Trees?

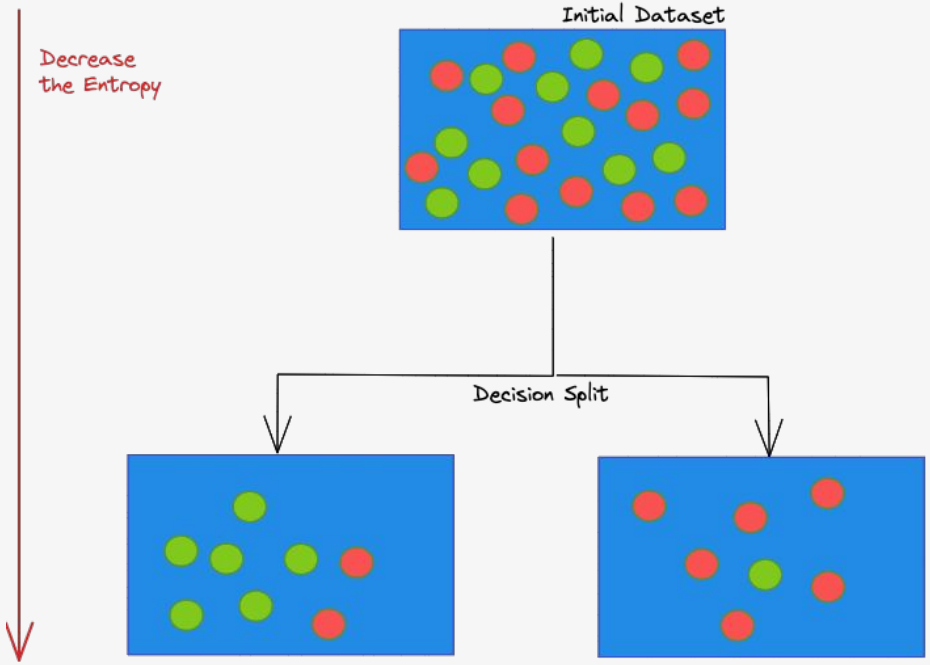

- The goal is to tidy the data.
- You try to separate your data and group the samples together in the classes they belong to.
- You maximize the purity of the groups as much as possible each time you create a new node of the tree
- Of course, at the end of the tree, you want to have a clear answer.

Mathematical Definition of Entropy

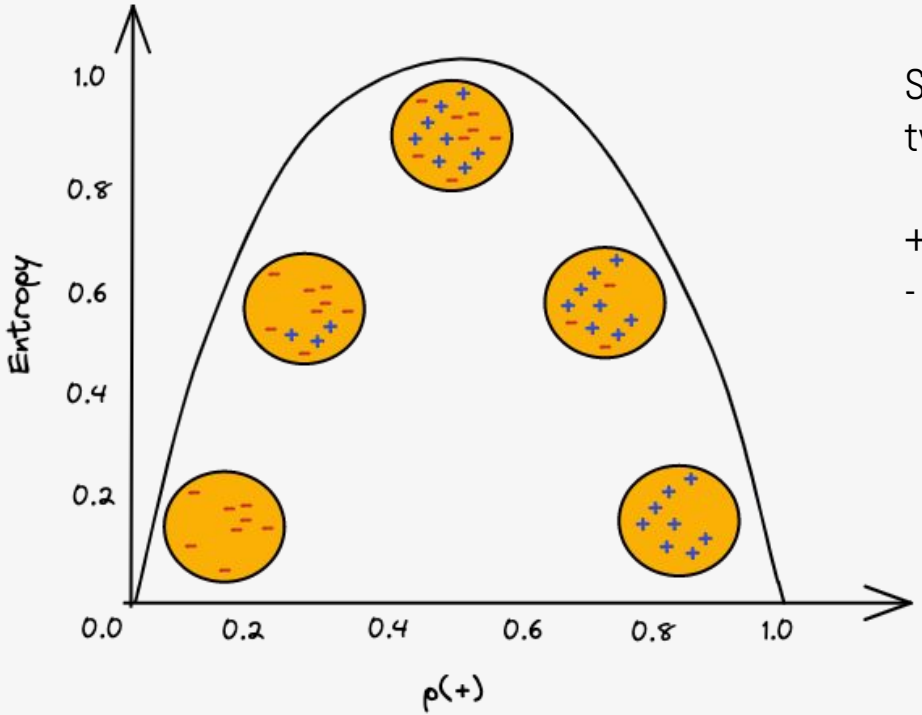

Suppose a set of N items, these items fall into two categories:

\+ gain > 50k (k)
\- gain <= 50k (m)

$$p = \frac{k}{N}, q = \frac{m}{N}$$

$$Entropy = -p \log p - q \log q$$

Generalization

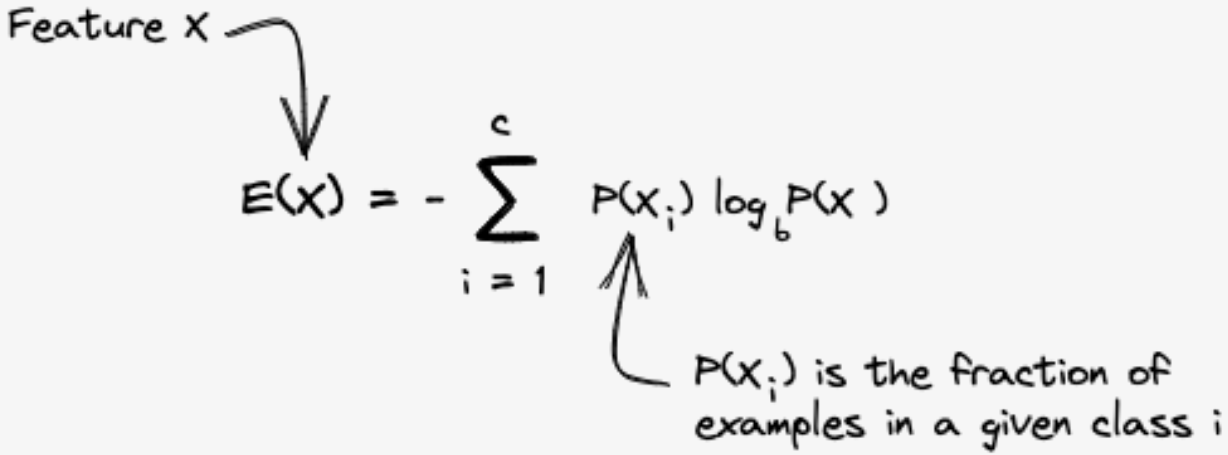

<= 50k.	17288
> 50k.	5487

```
from scipy.stats import entropy
entropy(df_train.high_income.value_counts(), base=2)
0.7965702796015677
```

Entropy using the frequency table of two attributes

		High Income		
		<= 50k	> 50k	
Age	<=37	7206	3883	11089 (48%)
	>37	10082	1604	11686 (52%)

$$E(T|X) = \sum_{c \in X} \frac{|X_c|}{|X|} E(T|X_c)$$

```
cross = pd.crosstab(
          df_train.age <= df_train.age.median(),
          df_train.high_income)
```

```
0.486894 * entropy(cross.iloc[0], base=2) \
+ 0.513106 * entropy(cross.iloc[1], base=2)
0.7509335429830957
```

Information Gain

$$IG(T,X) = E(T) - E(T|X)$$

Information Gain from X on T

The information gain is based on the **decrease in entropy after a dataset is split** on an attribute.

Constructing a decision tree is all about finding attribute that returns the **highest information gain** (i.e., the most homogeneous branches).

$$Gini(x) = 1 - \sum_{i=1}^{c} P(x_i)^2$$

$$Entropy(x) = - \sum_{i=1}^{c} P(x_i) \log_b P(x_i)$$

Gini index or Entropy is the criterion for calculating **Information Gain**. Both of them are measures of impurity of a node.

```
from sklearn.tree import plot_tree
```

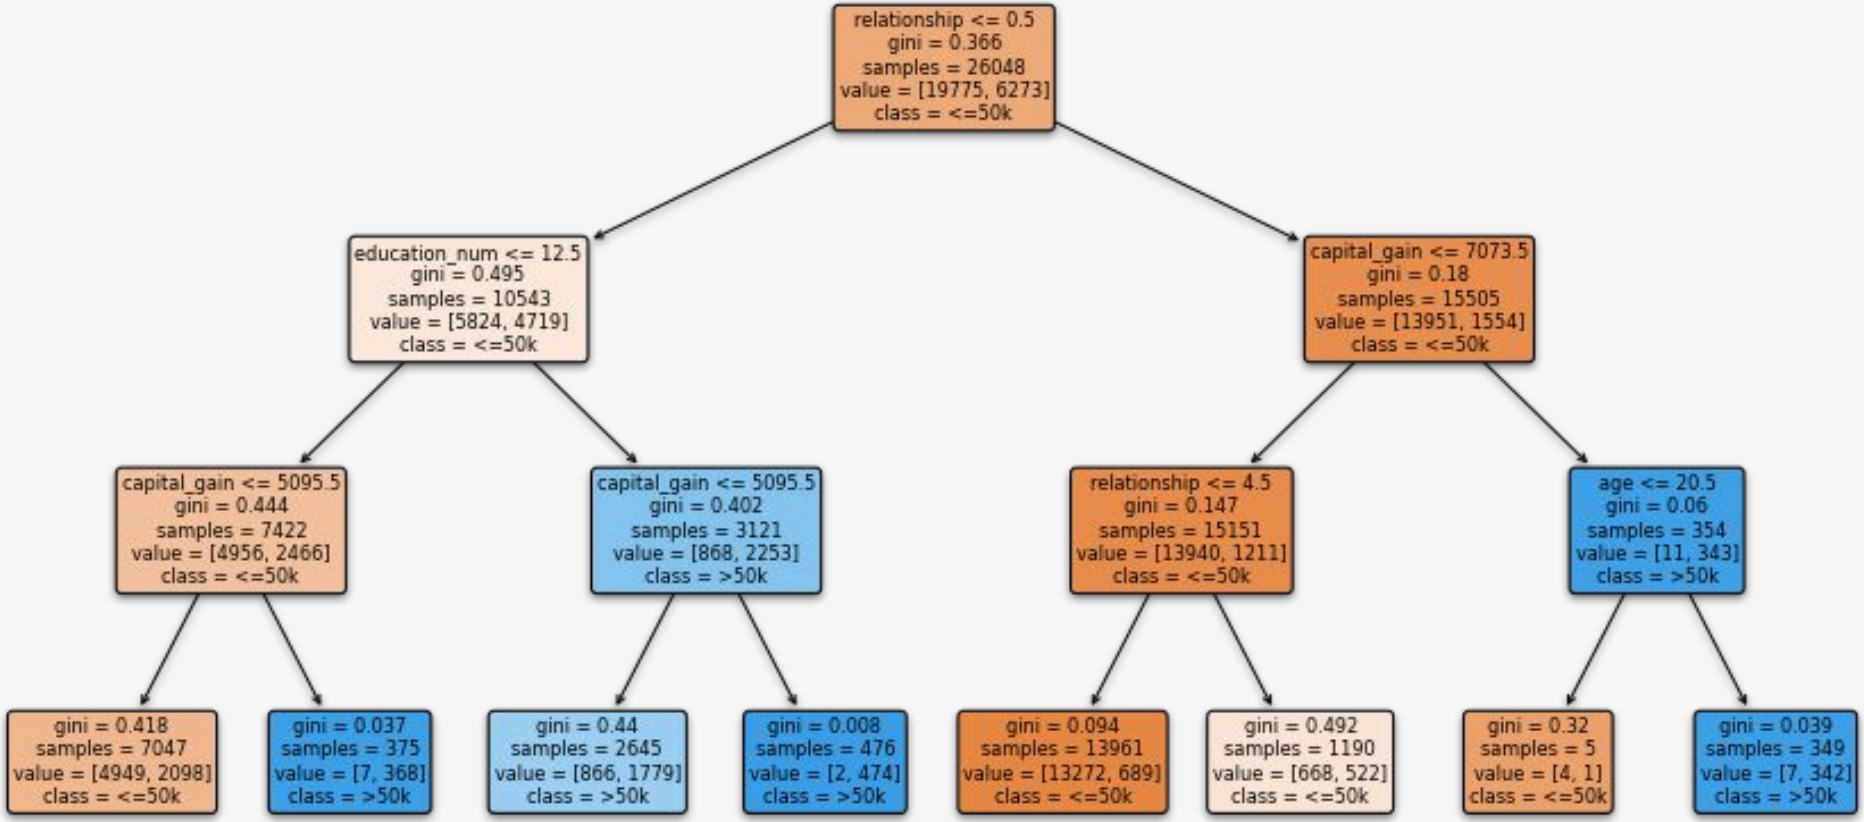

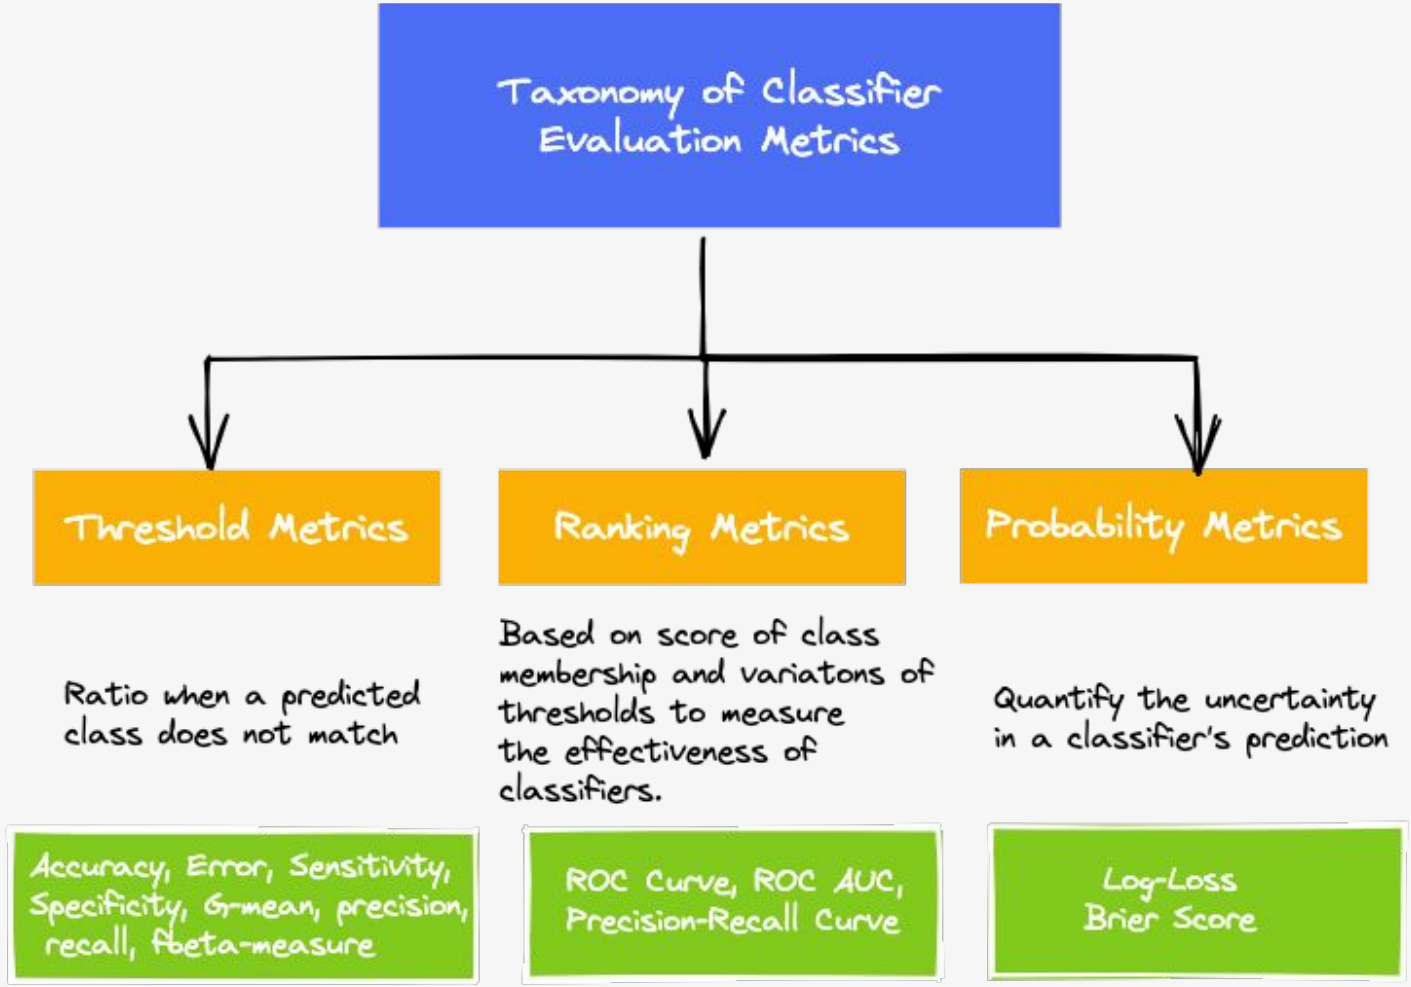
Taxonomy of Classifier Evaluation Metrics
Threshold Metrics
Ranking Metrics
Probability Metrics
Ratio when a predicted class does not match
Based on score of class membership and variatons of thresholds to measure the effectiveness of classifiers.
Quantify the uncertainty in a classifier's prediction
Accuracy, Error, Sensitivity, Specificity, G-mean, precision, recall, fbeta-measure
ROC Curve, ROC AUC, Precision-Recall Curve
Log-Loss
Brier Score

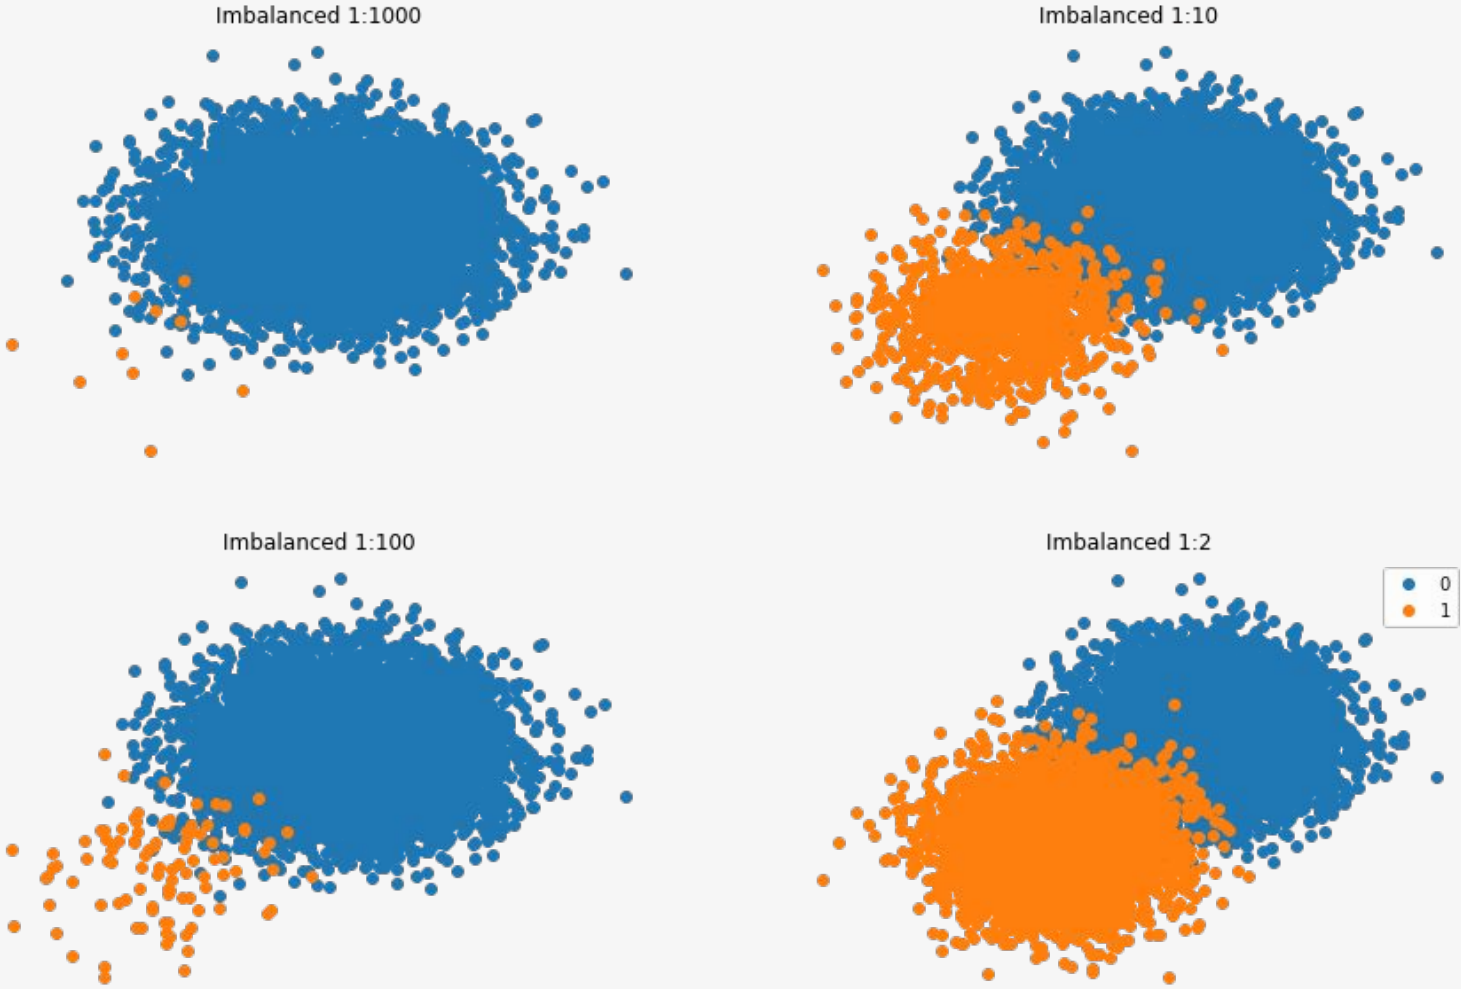
Imbalanced 1:1000
Imbalanced 1:10
Imbalanced 1:100
Imbalanced 1:2
0
1

How to choose an evaluation metric?

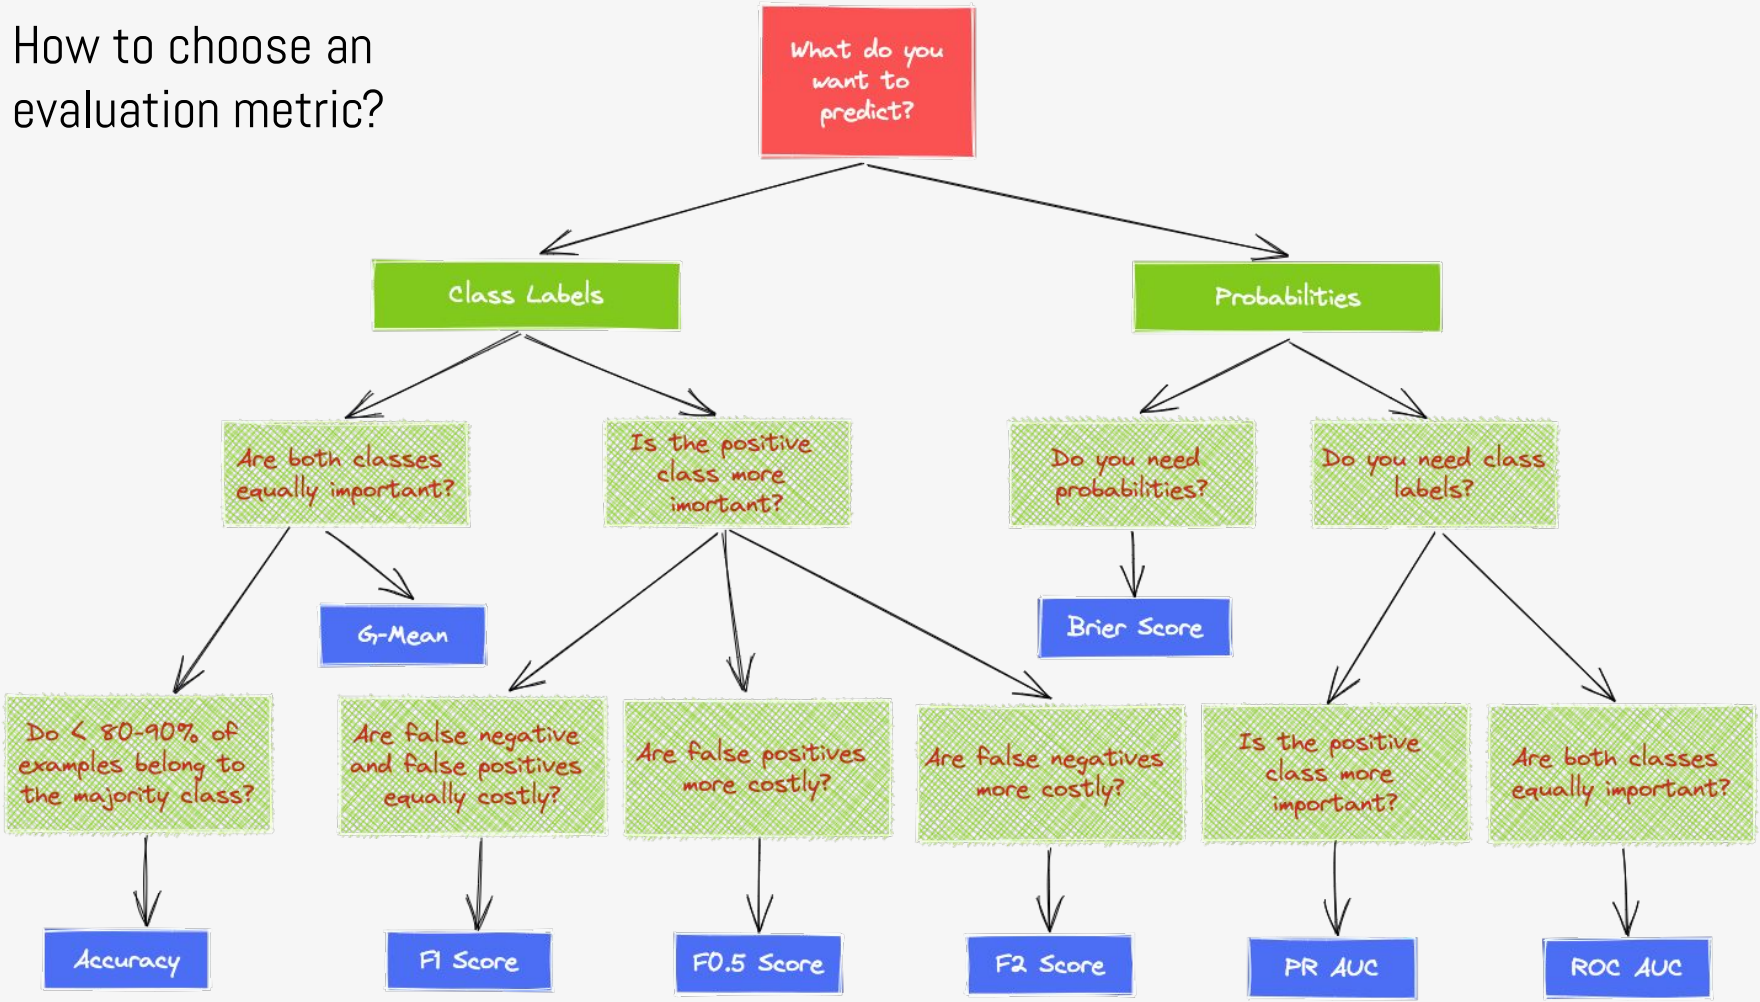

@Brownlee, Jason. Imbalanced classification with python.

Confusion Matrix

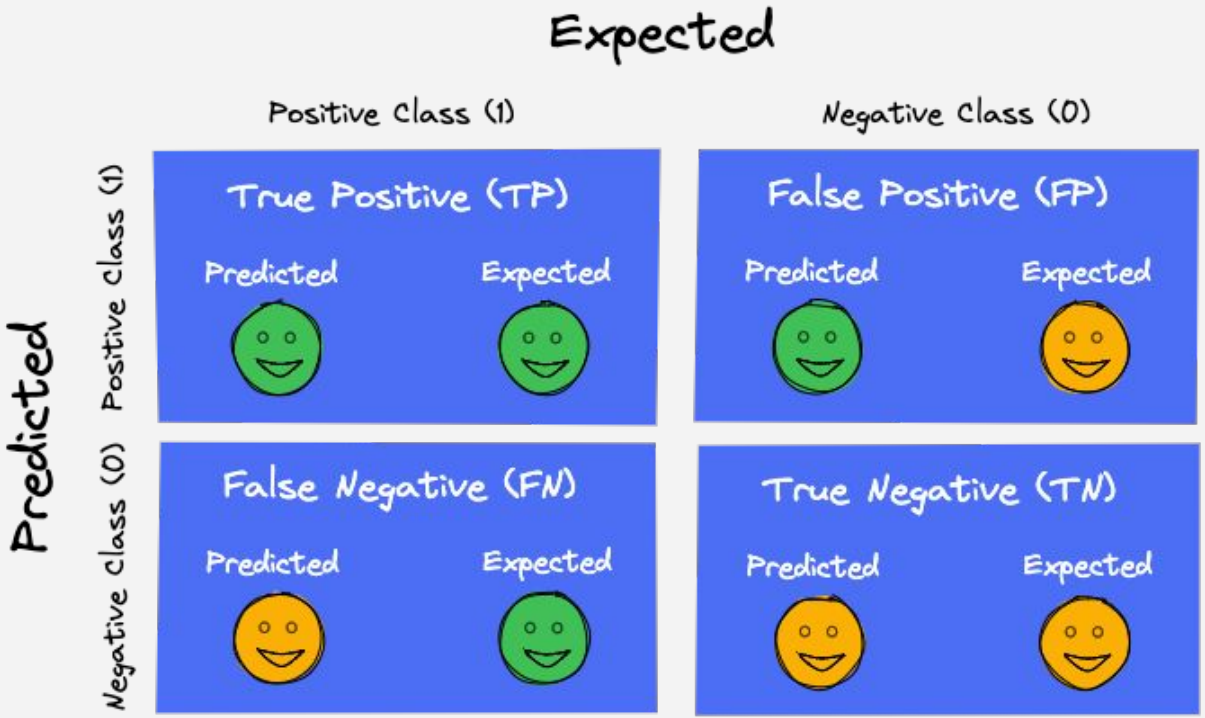

$$\text{Accuracy} = \frac{TP + TN}{TP + FN + FP + TN}$$

$$\text{Error} = 1 - \text{Accuracy}$$

Confusion Matrix

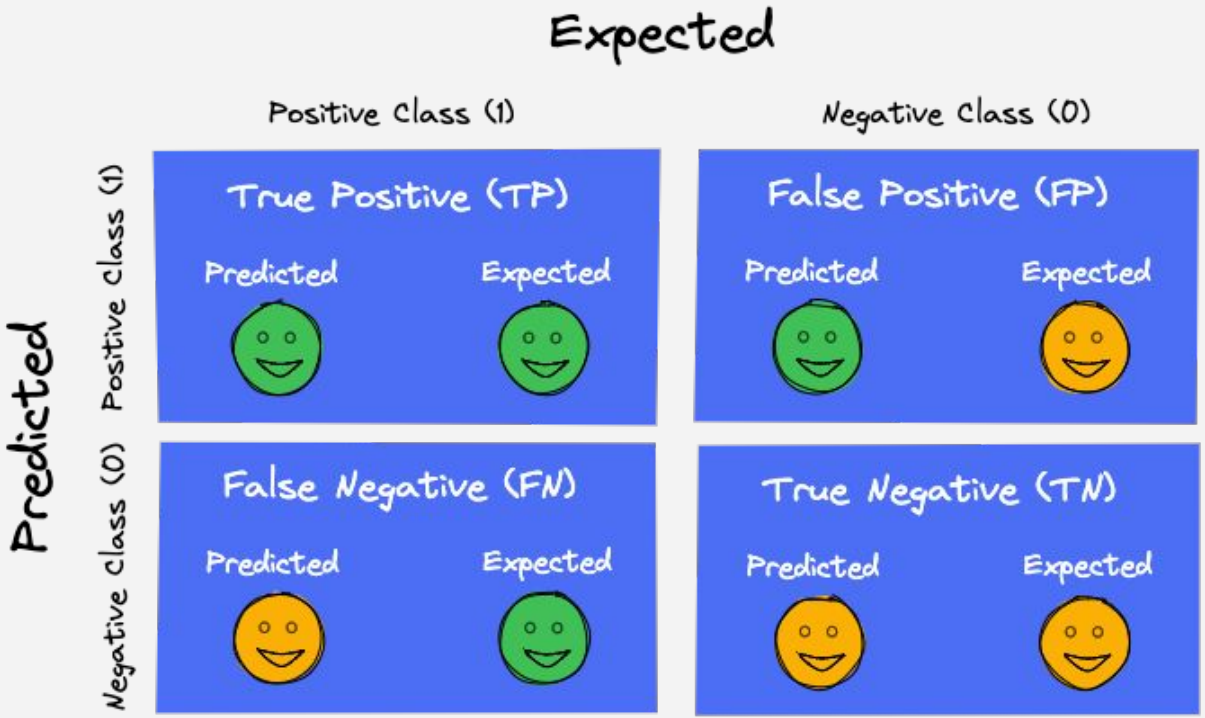

$$\text{Sensitivity} = \frac{TP}{TP + FN}$$

$$\text{Specificity} = \frac{TN}{FP + TN}$$

$$\text{G-mean} = \sqrt{\text{Sensitivity} \times \text{Specificity}}$$

Confusion Matrix

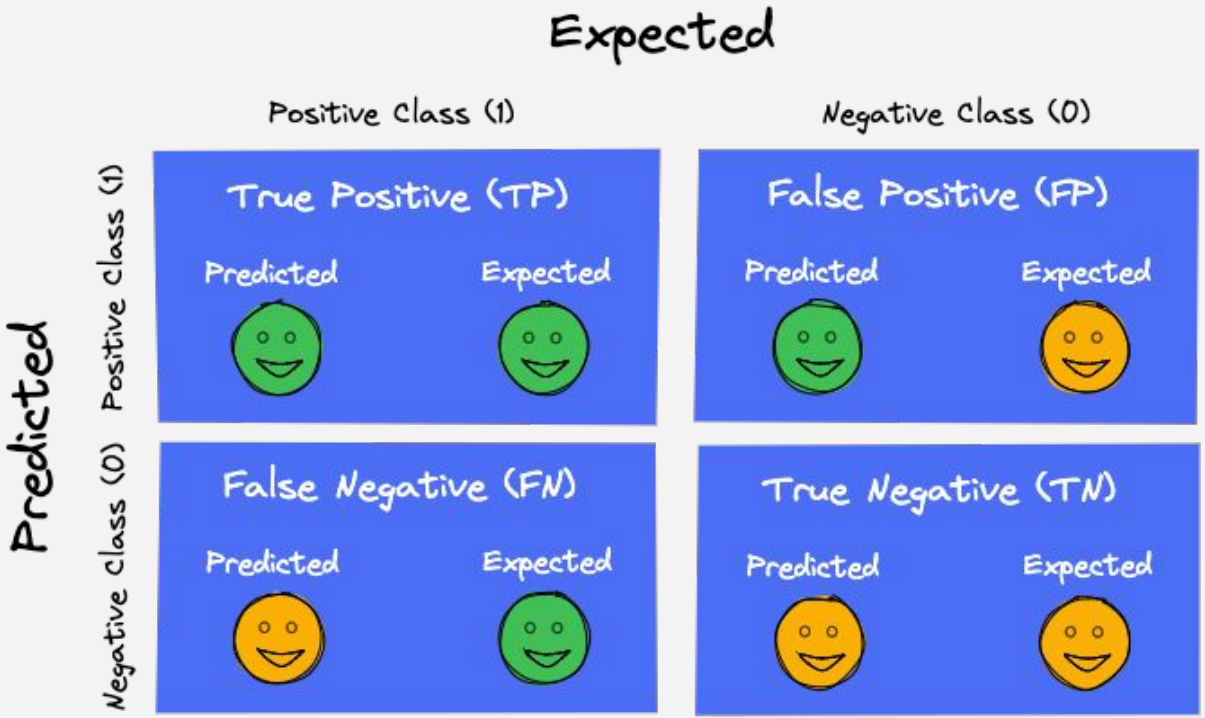

$$\text{Precision (positive predicte value - PPV)} = \frac{TP}{TP + FP}$$

$$\text{Precision (negative predicte value - NPV)} = \frac{TN}{TN + FN}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$

Confusion Matrix

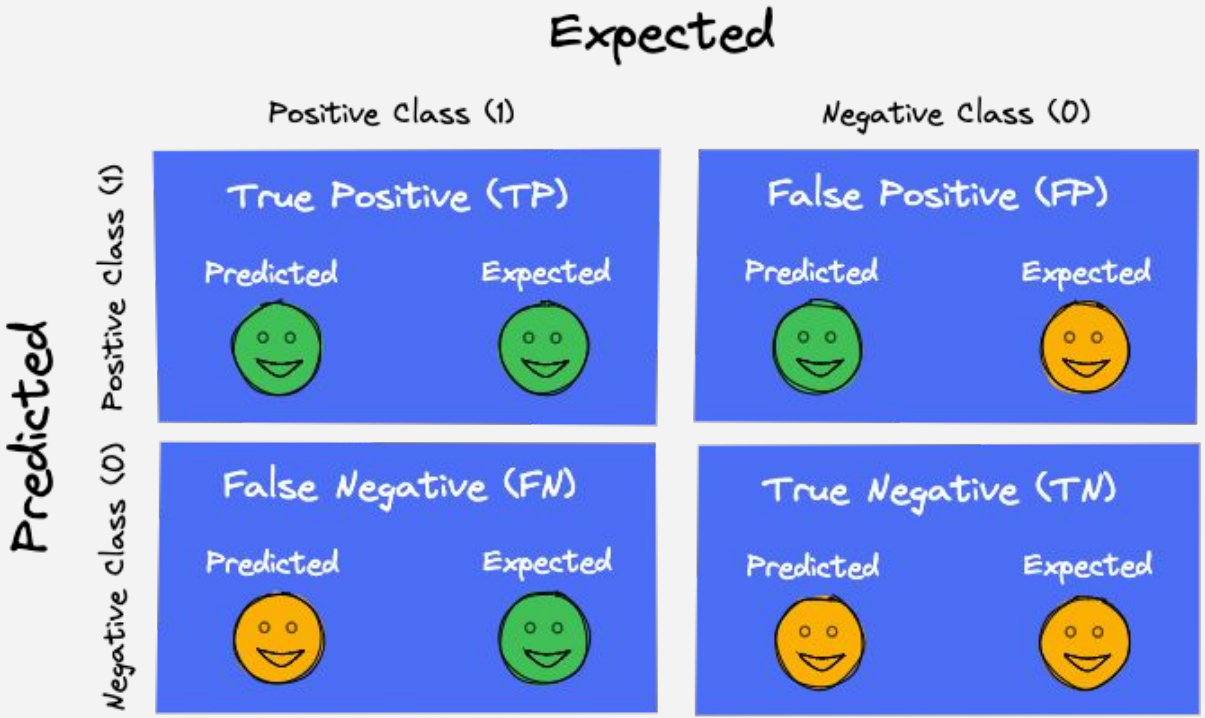

$$\text{Fbeta-measure} = \frac{(1 + \beta^2) \times \text{Precision} \times \text{Recall}}{\beta^2 \times \text{Precision} + \text{Recall}}$$

$$\beta == \begin{cases} 0.5, \text{ more weight on precision} \\ 1.0, \text{ balance on weight PR and RE} \\ 2.0, \text{ less weight on precision} \end{cases}$$

Rank metrics are more concerned with evaluating classifiers based on **how effective** they are at separating classes.

These metrics require that a **classifier predicts a score** or a probability of class membership. From this score, **different thresholds** can be applied to **test the effectiveness of classifiers**. Those models that maintain a good score across a range of thresholds will have good class separation and will be ranked higher.

Receiver Operating Characteristic (ROC)

$$TPR = \frac{TP}{TP + FN}$$

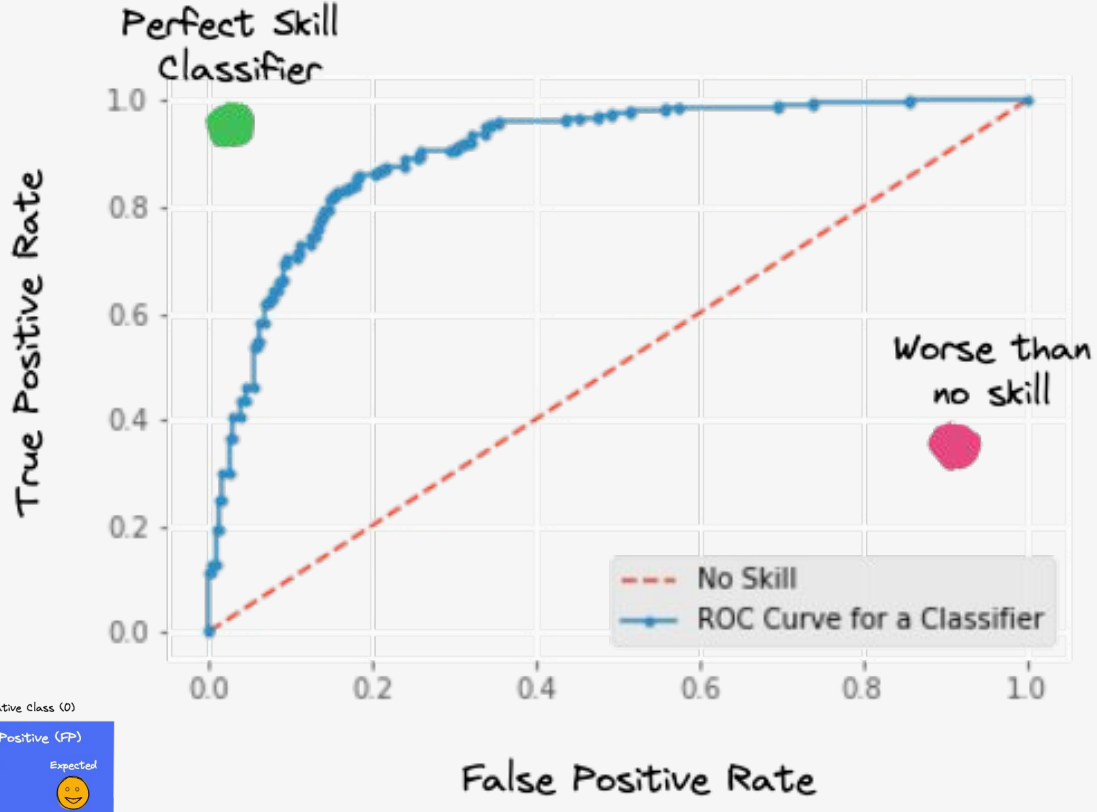

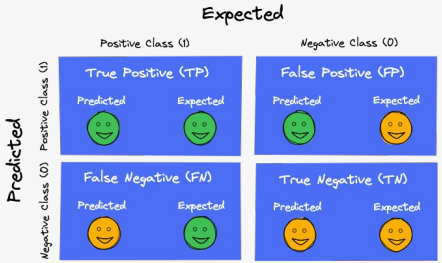

$$FPR = \frac{FP}{FP + TN}$$

Precision-Recall (PR) Curve

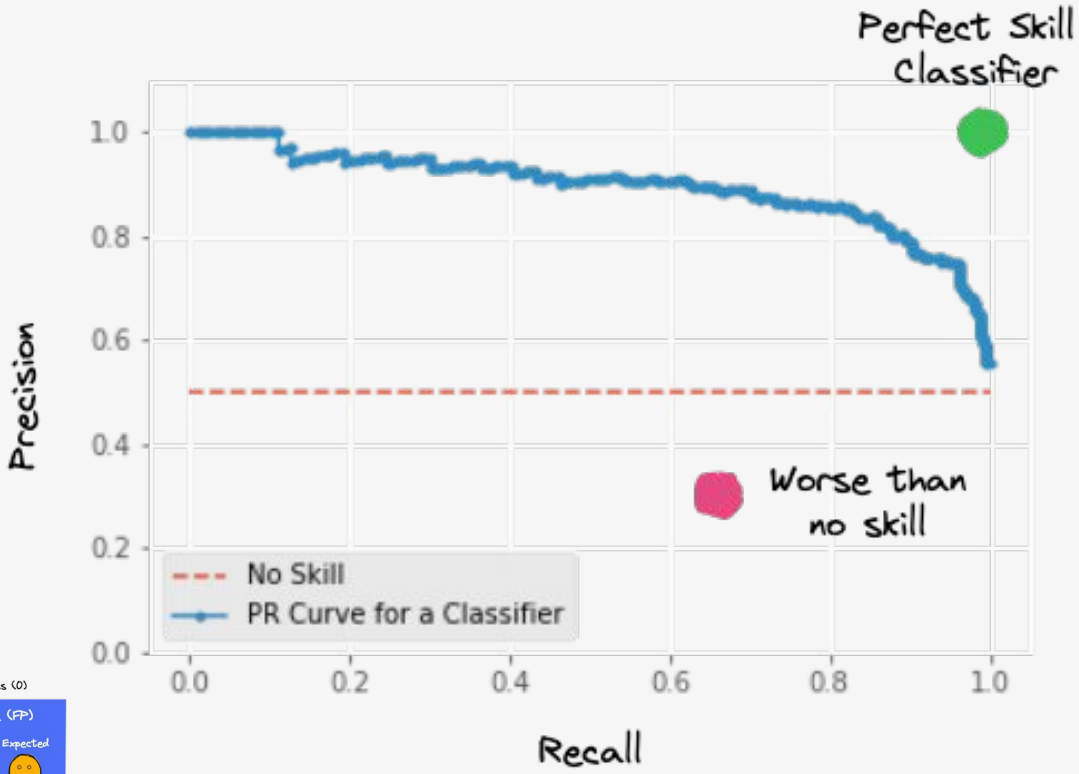

$$\text{Precision} = \frac{TP}{TP + FP}$$

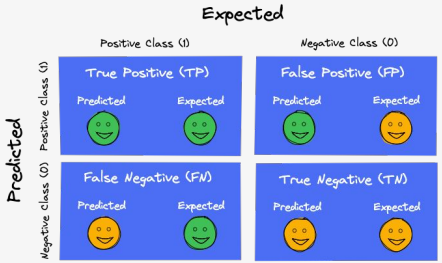

$$\text{Recall} = \frac{TP}{TP + FN}$$

Hands ON

Python

learn

W&B

Jupyter

CO

Streamlit

ANACONDA

CONDA

Optional

matplotlib

pandas

seaborn

plotly

UCI Machine Learning Repository

Center for Machine Learning and Intelligent Systems

Adult Data Set

Download: Data Folder, Data Set Description

Abstract: Predict whether income exceeds $50K/yr based on census data. Also known as "Census Income" dataset.

Data Set Characteristics:	Multivariate	**Number of Instances:**	48842	**Area:**	Social
Attribute Characteristics:	Categorical, Integer	**Number of Attributes:**	14	**Date Donated**	1996-05-01
Associated Tasks:	Classification	**Missing Values?**	Yes	**Number of Web Hits:**	2437279

Source:

Donor:

Ronny Kohavi and Barry Becker
Data Mining and Visualization
Silicon Graphics.
e-mail: ronnyk '@' live.com for questions.

Exploratory Data Analysis

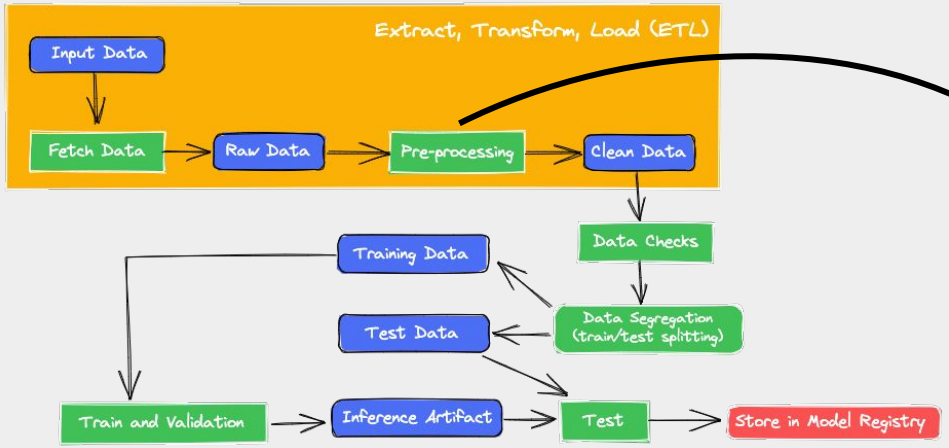

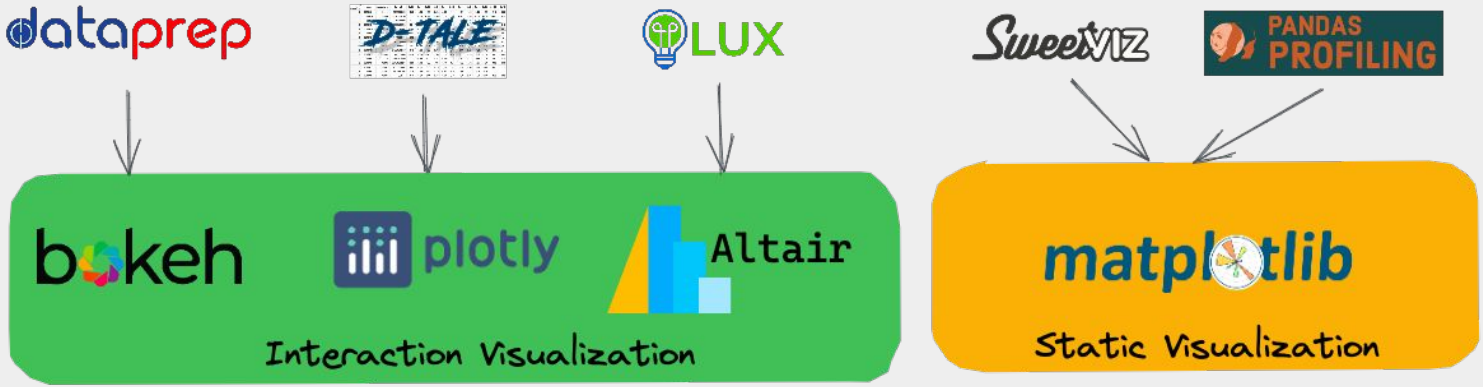

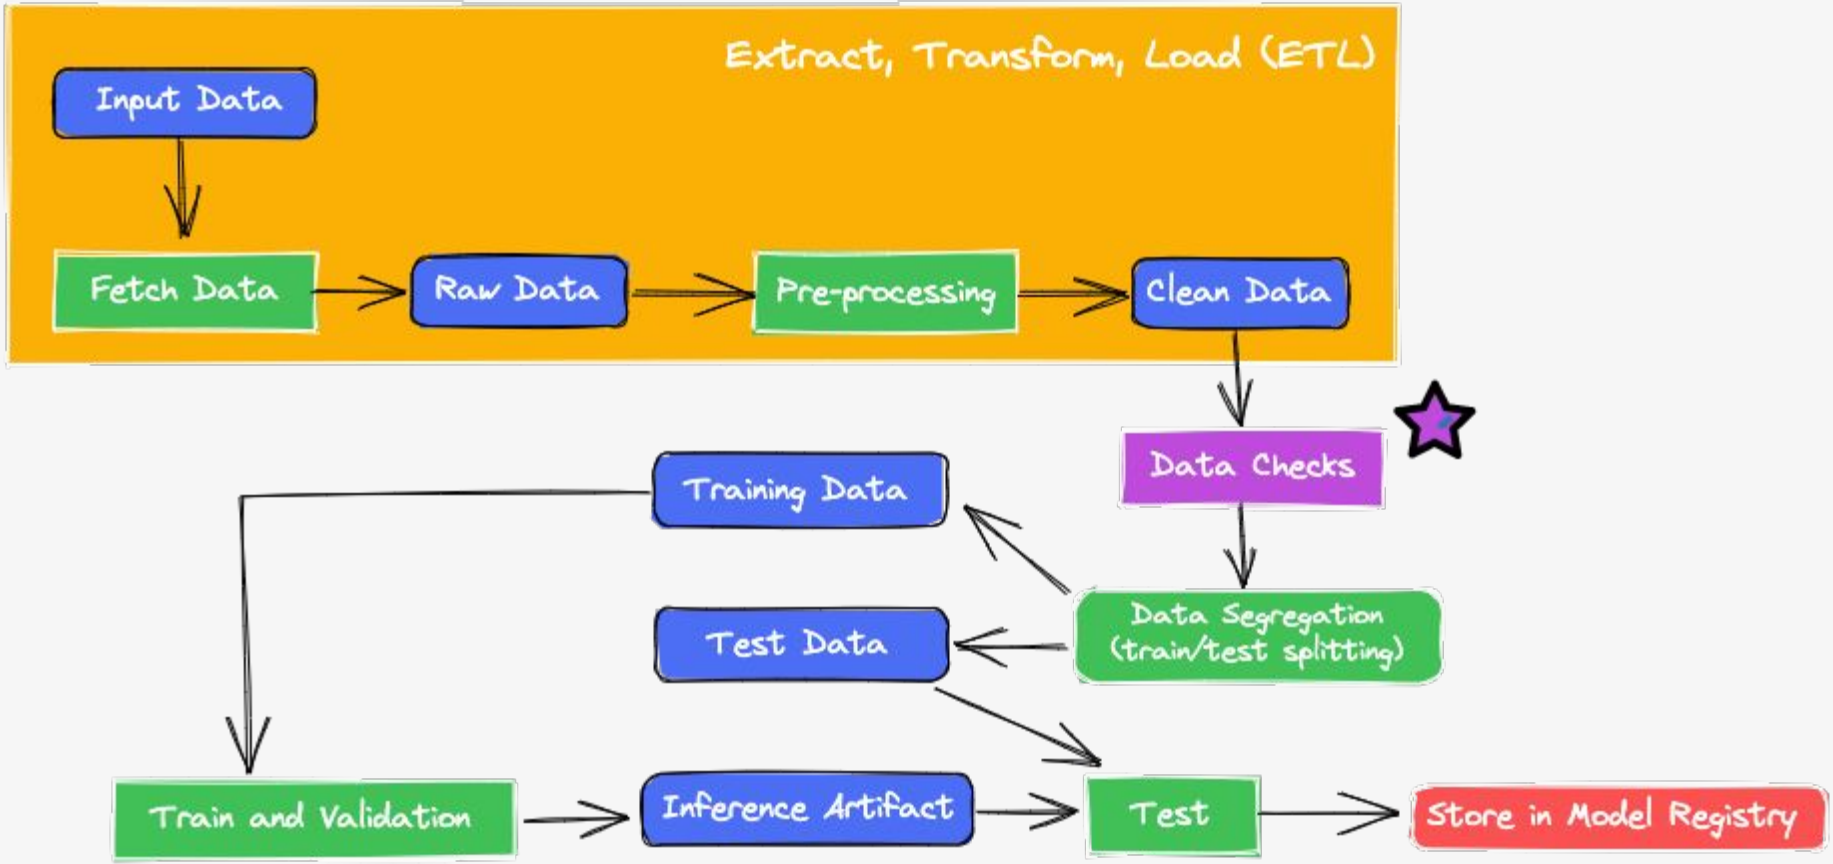
Extract, Transform, Load (ETL)
Input Data
Fetch Data
Raw Data
Pre-processing
Clean Data
Data Checks
Training Data
Test Data
Data Segregation (train/test splitting)
Train and Validation
Inference Artifact
Test
Store in Model Registry

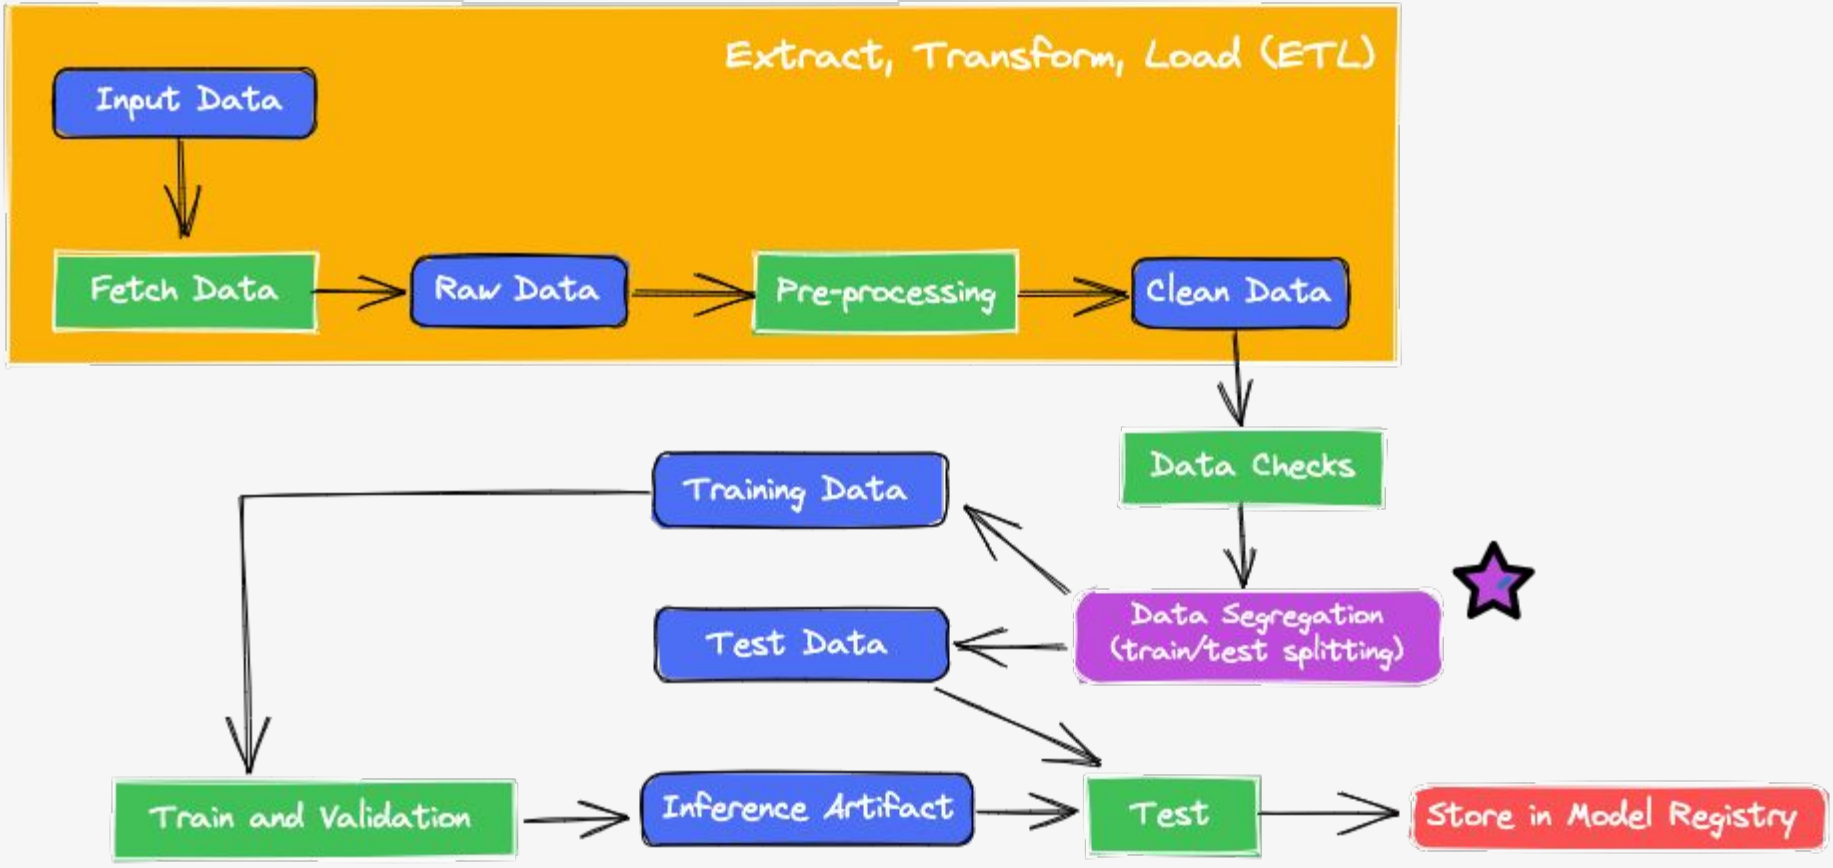
Extract, Transform, Load (ETL)
Input Data
Fetch Data
Raw Data
Pre-processing
Clean Data
Data Checks
Training Data
Test Data
Data Segregation
(train/test splitting)
Train and Validation
Inference Artifact
Test
Store in Model Registry

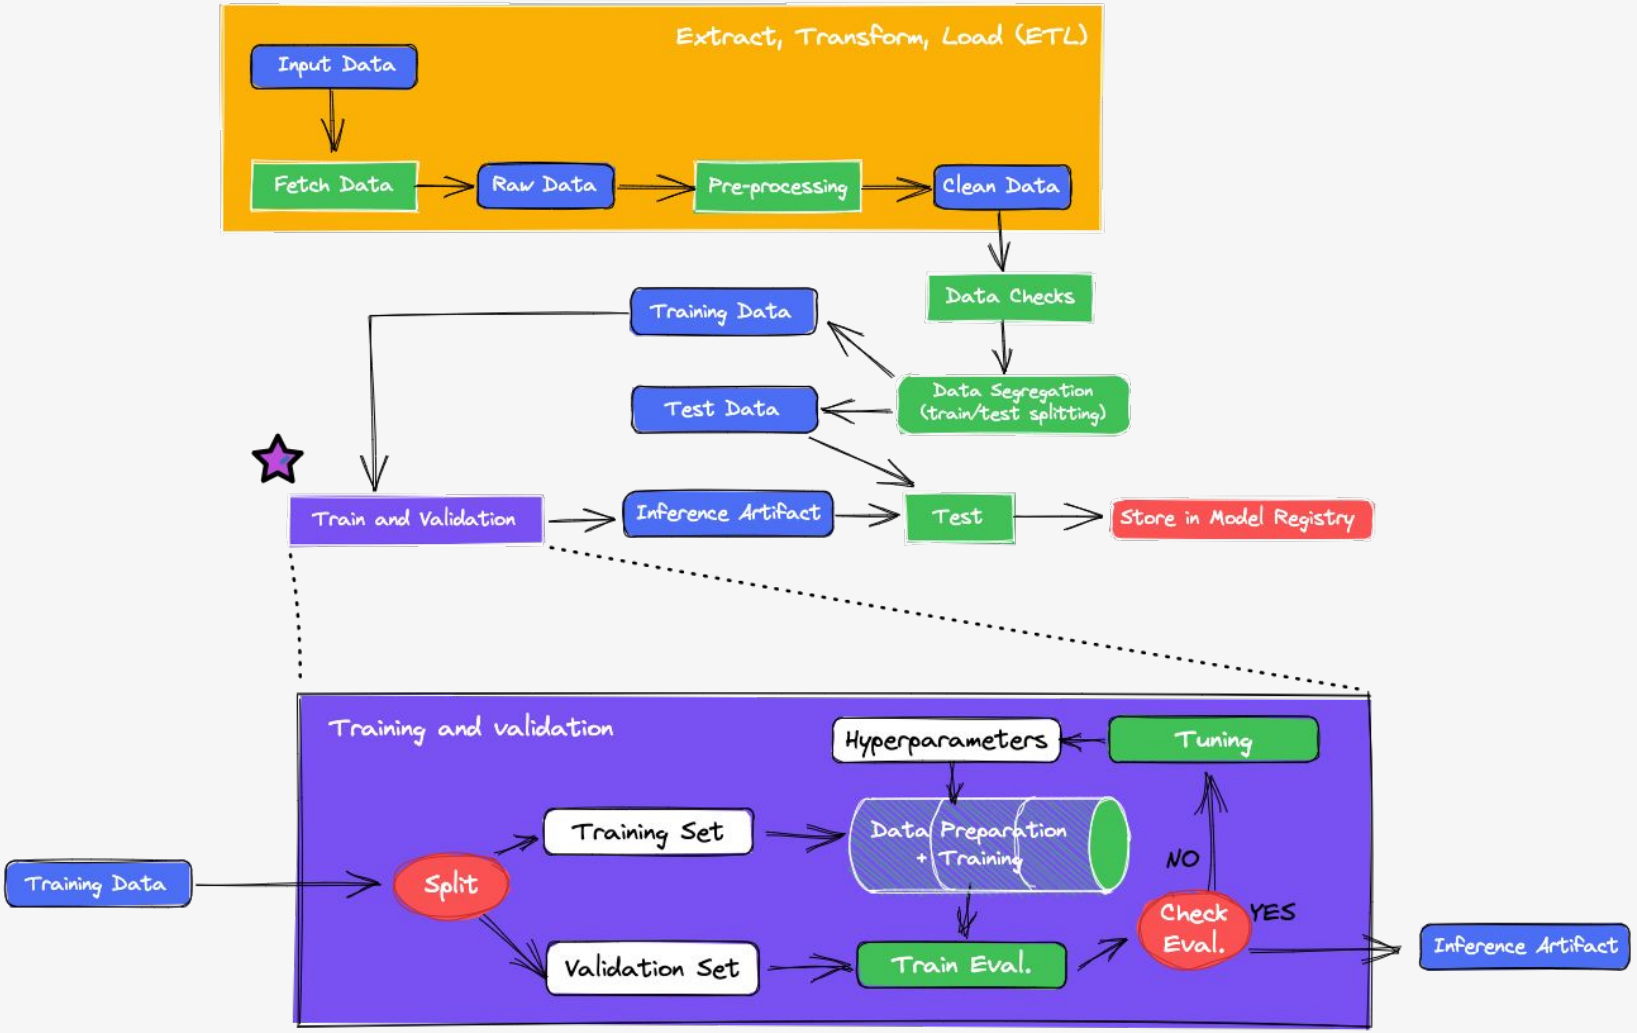
Extract, Transform, Load (ETL)
Input Data
Fetch Data
Raw Data
Pre-processing
Clean Data
Data Checks
Training Data
Test Data
Data Segregation (train/test splitting)
Train and Validation
Inference Artifact
Test
Store in Model Registry
Training and validation
Hyperparameters
Tuning
Training Data
Split
Training Set
Validation Set
Data Preparation + Training
Train Eval.
Check Eval.
NO
YES
Inference Artifact


Raw Data
Data Segregation
Train Set
Test Set
Holdout
Train Set
Validation Set
cross-validation
Split

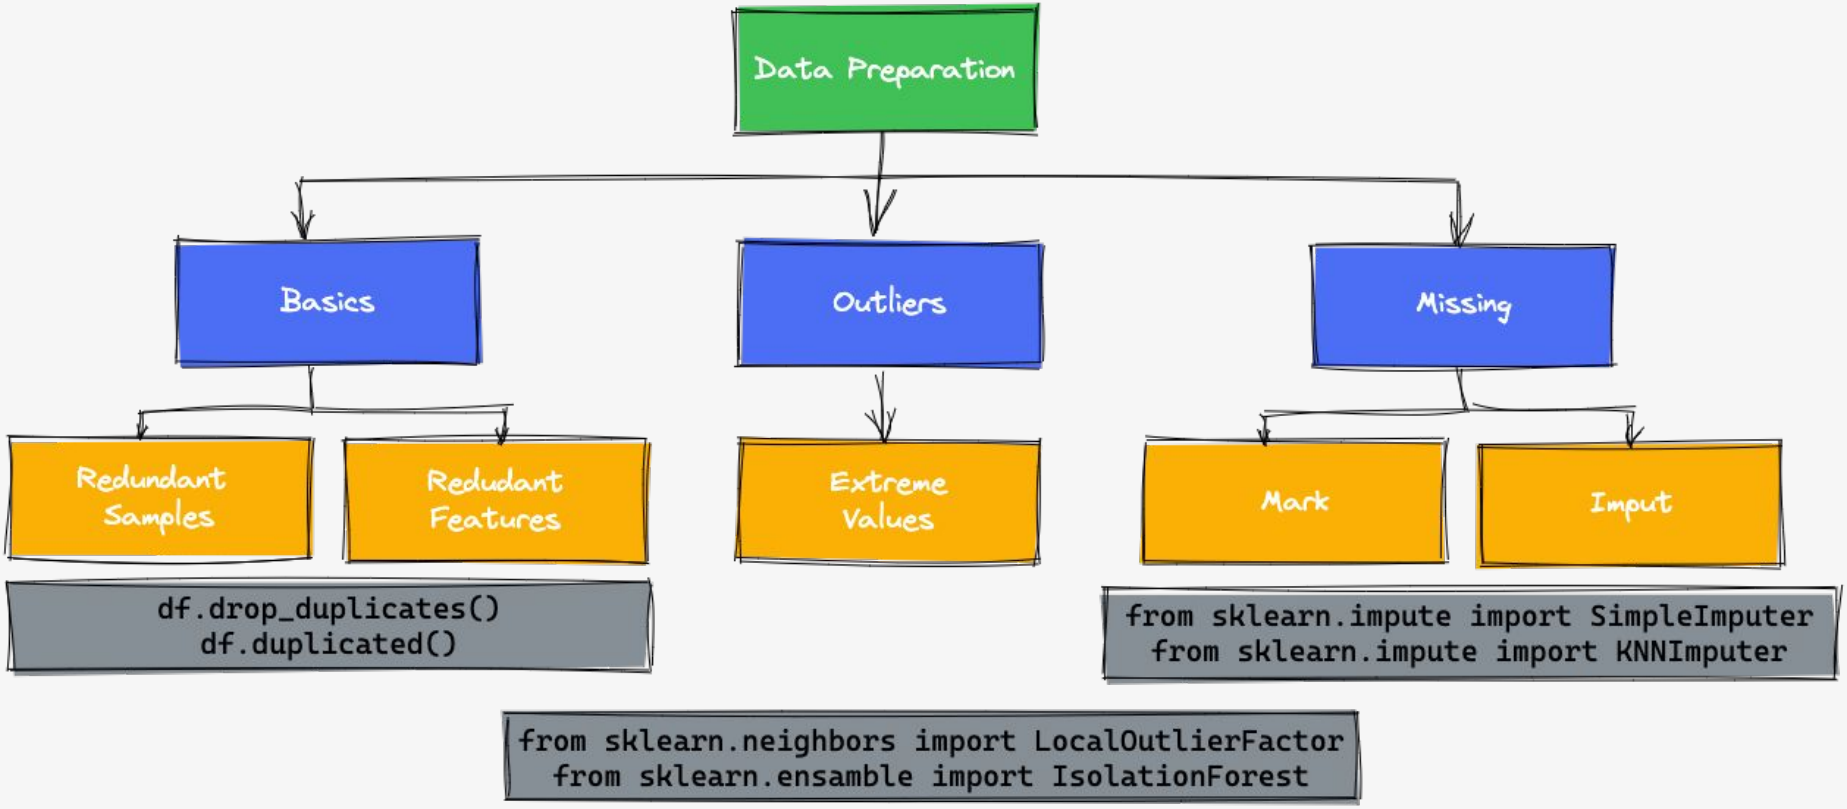
Data Preparation
Basics
Outliers
Missing
Redundant Samples
Redudant Features
Extreme Values
Mark
Imput
df.drop_duplicates()
df.duplicated()
from sklearn.impute import SimpleImputer
from sklearn.impute import KNNImputer
from sklearn.neighbors import LocalOutlierFactor
from sklearn.ensamble import IsolationForest

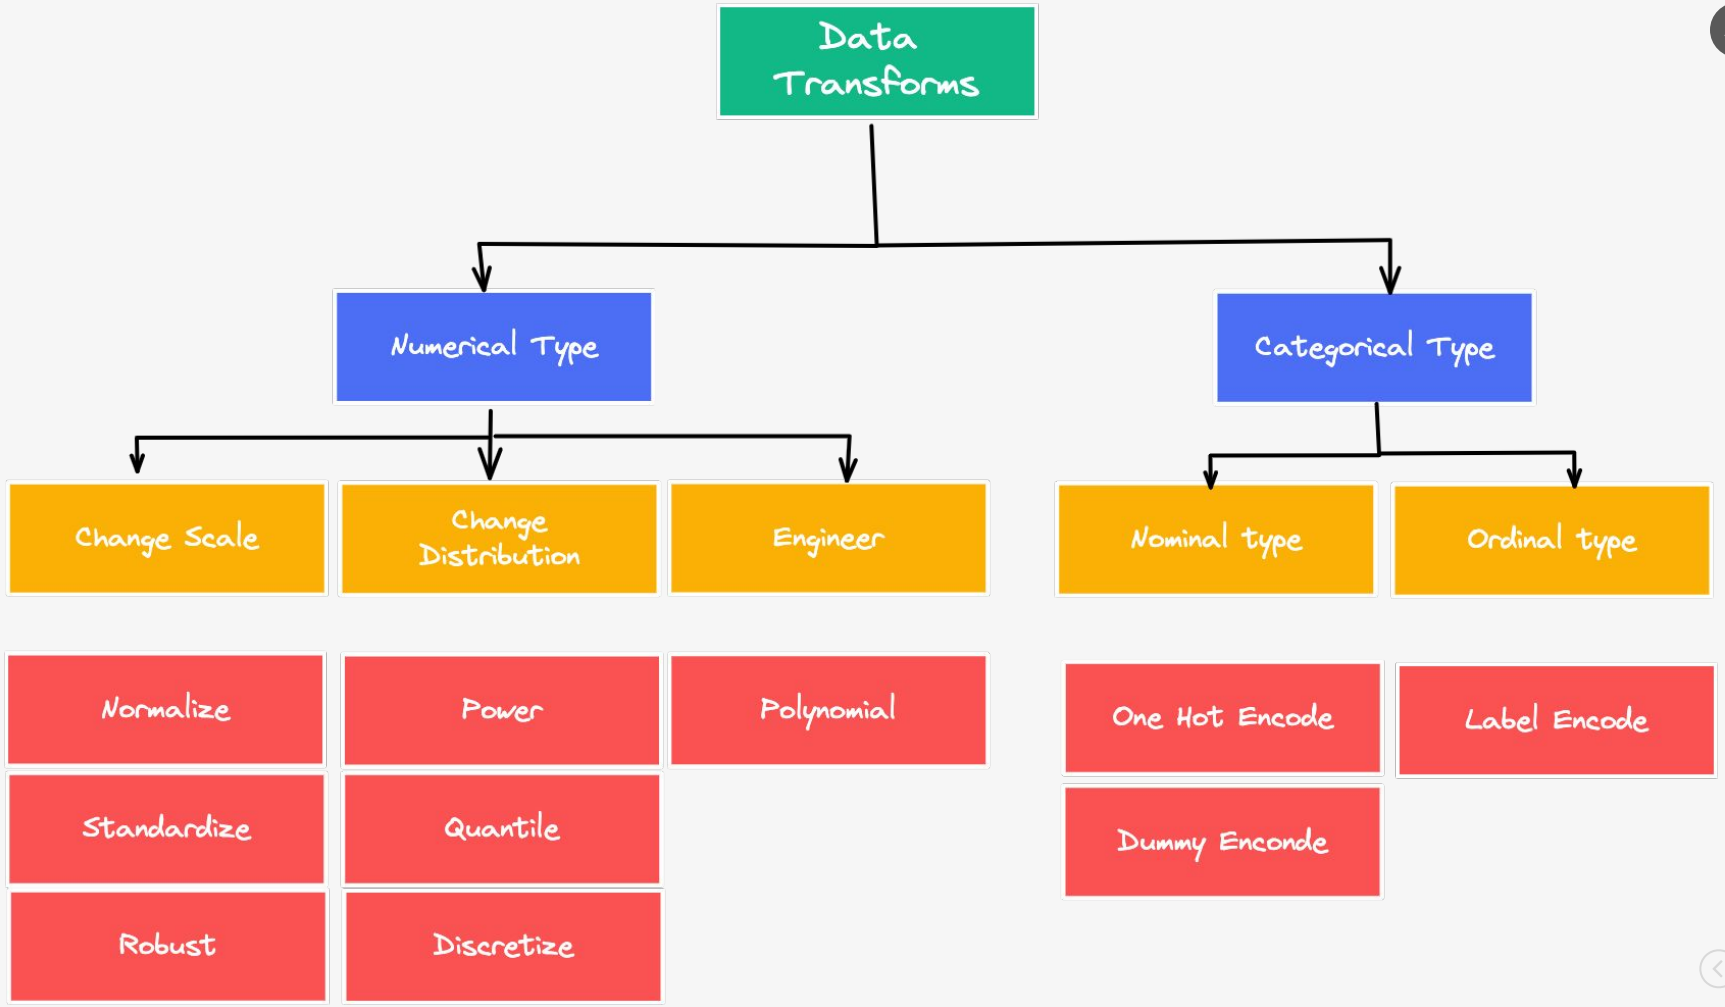
Data Transforms
Numerical Type
Categorical Type
Change Scale
Change Distribution
Engineer
Nominal type
Ordinal type
Normalize
Power
Polynomial
One Hot Encode
Label Encode
Standardize
Quantile
Dummy Enconde
Robust
Discretize

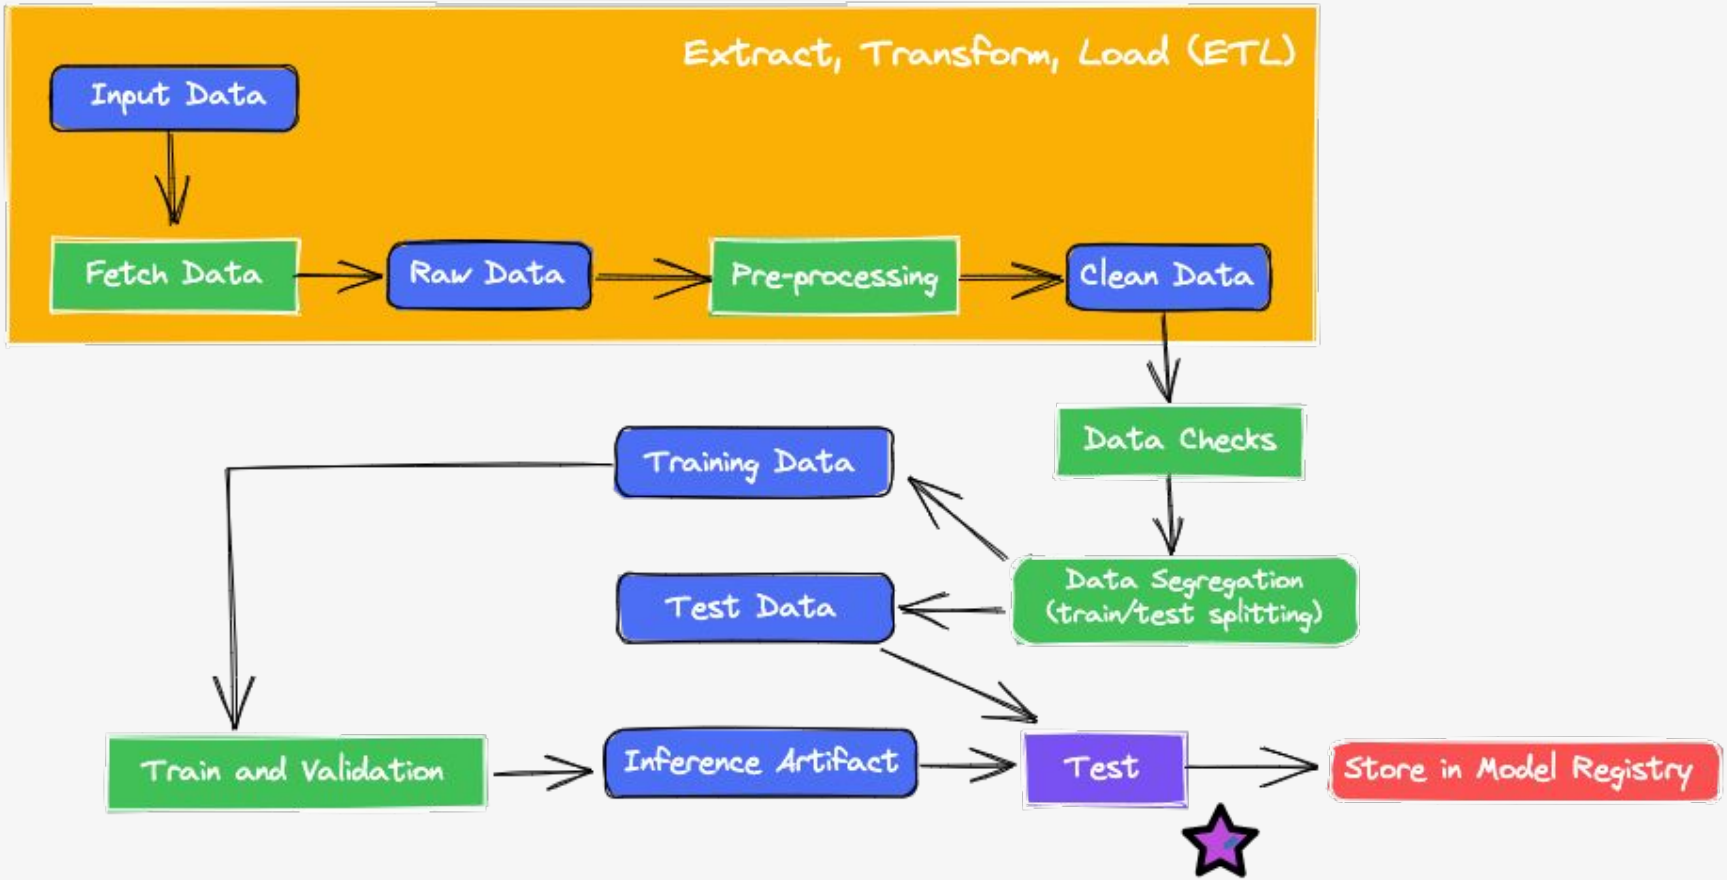
Extract, Transform, Load (ETL)
Input Data
Fetch Data
Raw Data
Pre-processing
Clean Data
Data Checks
Training Data
Test Data
Data Segregation (train/test splitting)
Train and Validation
Inference Artifact
Test
Store in Model Registry